

Efficient Verification of Pairing Computation: Schwartz-Zippel Lemma for Miller loop

Shramee Srivastav
MIST.cash — FOCBB
Email: shramee@proton.me

September 7, 2025

Abstract

In this paper, we examine the advances in efficiently proving the correctness of elliptic curve pairing computations. Rather than focusing on raw pairing computation itself, our work centers on practical, sound, and high-performance techniques for proving that pairing operations have been performed correctly. These advances are crucial for making SNARK-based proof verification practical in computationally constrained environments—including BitcoinVM, Ethereum smart contracts, ZKVMs, and recursive zero-knowledge proof systems—where traditional pairing verification would be prohibitively expensive.

We systematically present a progression of state-of-the-art techniques: El Housni’s foundational R1CS pairing optimizations [Hou22], Feltroidprime’s enhanced extension field arithmetic [Fel23], and Novakovic and Eagen’s final exponentiation elimination [NE24].

Building upon these foundations, we introduce our novel contribution that extends the Schwartz-Zippel verification approach specifically tailored for Miller loop applications. Our technique, besides improving the Miller loop performance by 50% simultaneously reduces auxiliary data size for verification—a critical consideration for blockchain deployment where calldata directly affects transaction costs. We conclude with a comprehensive performance analysis of the current state-of-the-art pairing verification [Con] and the impact of our contributions.

Contents

1	Introduction	2
1.1	Elliptic curve pairings	2
1.2	Constraint-Based Computation Models	3
1.3	Related Work	3
1.4	Our Contribution	3
1.5	Our Approach	4
2	Background and Technical Roadmap	4
2.1	Pairing Computation	4
2.2	Schwartz–Zippel lemma	5
2.3	Roadmap	5

3	Pairings in Rank-1 Constraint Systems	5
3.1	Cost of inversions	6
3.2	Miller loop	6
3.2.1	Affine coordinates	6
3.2.2	Avoid doubling	6
3.2.3	Line evaluations	6
3.3	Algebraic Torii	7
3.4	Final exponentiation	7
3.5	Performance	7
4	Faster Extension Field Multiplications	7
4.1	Isomorphisms among towered to direct extensions	7
4.2	Multiplication and Schwartz–Zippel lemma	8
4.3	Deferred Evaluation and Batching	8
4.4	Performance	9
5	Eliminating the Final Exponentiation	9
5.1	Equivalence class	9
5.2	Equivalence via Residue Check	10
5.3	Finding optimal parameters	10
5.3.1	Hasse bound and exponentiation by r	10
5.3.2	Parameters for BN254 curve	10
6	Miller loop Step as Polynomial	11
6.1	Core Innovation	11
6.1.1	Comparison setup	11
6.1.2	Previous scheme	12
6.1.3	Our scheme	12
6.2	Bit-Type Polynomials for 3-pairing	13
6.2.1	Zero Bit Equation	13
6.2.2	Non-Zero Bit Equation	13
6.2.3	Correction Step Equation	14
6.3	Polynomial Accumulation and implementation	14
6.3.1	Prover	15
6.3.2	Verifier	15
6.4	Performance	15

1 Introduction

1.1 Elliptic curve pairings

Elliptic curve pairings represent fundamental building blocks in modern cryptographic protocols, serving as the mathematical foundation for numerous cryptographic systems including Groth16 [Gro16], PLONK [GWC19], BLS signature schemes [BGW⁺22], and KZG polynomial commitment schemes [KZG10]. At their core, elliptic curve pairings provide a bilinear map [ABLR14] $e : \mathbb{G}_1 \times \mathbb{G}_2 \rightarrow \mathbb{G}_T$ that enables a form of an encrypted multiplication over elliptic curve groups. This bilinearity property — where $e(aP, Q) = e(P, aQ) = e(P, Q)^a$ for scalar a and points $P \in \mathbb{G}_1, Q \in \mathbb{G}_2$ — has enabled breakthrough

advances across cryptographic applications, from identity-based encryption to the zero-knowledge systems securing modern blockchain infrastructure.

1.2 Constraint-Based Computation Models

In the context of arithmetization protocols and virtual machines (such as Starkware’s Cairo VM [Sta22], R1CS [GGPR13], Plonkish [GWC19], etc.), where field operations are expressed as algebraic constraints, the computational landscape differs dramatically from native implementations. Operations that are traditionally expensive can become remarkably efficient. Field inversions, for example, can be computed at a fraction of their usual cost, often reducing their relative costs from over $100\times$ to about the same cost as multiplication. Conversely, seemingly straightforward operations may require complex constraint representations.

Recognizing and exploiting these unique computational properties is essential for unlocking hidden performance potential. Youssef El Housni’s pioneering work **Pairings in Rank-1 Constraint Systems** [Hou22] represents, to our knowledge, the first systematic exploration of these optimization opportunities, providing a breakthrough in the practicality of recursive proving by thoughtfully leveraging the distinctive cost model inherent to constraint-based systems.

1.3 Related Work

The landscape of efficient pairing verification in constraint systems has evolved through several key contributions. El Housni’s **Pairings in Rank-1 Constraint Systems** [Hou22] established the theoretical framework for optimizing pairings in Rank-1 Constraint Systems, introducing techniques for efficient Miller loop implementation and final exponentiation that exploit the unique properties of constraint-based computation models. This work demonstrated that operations traditionally considered expensive in native implementations could be restructured to achieve remarkable efficiency gains in arithmetized environments.

Feltroidprime’s subsequent research on **enhanced extension field multiplication techniques** [Fel23] introduced polynomial verification methods using the Schwartz-Zippel lemma for $\mathbb{F}_{q^{12}}$ arithmetic operations. By representing extension field elements as polynomials and verifying their operations through polynomial arithmetic, this approach achieved significant performance improvements in the pairing computations, particularly for emulated pairing circuits where traditional field operations are costly.

Novakovic and Eagen’s work **On Proving Pairings** [NE24] presented breakthrough optimization strategies that enable the elimination of final exponentiation overhead through residue witness techniques. Their approach demonstrates that two elements in $\mathbb{F}_{q^{12}}$ can be proven equivalent without expensive exponentiation by introducing witness elements that satisfy specific algebraic relationships, effectively embedding the final exponentiation verification into the Miller loop computation itself.

1.4 Our Contribution

Building upon Feltroidprime’s foundational work on extension field arithmetic using polynomial verification [Fel23], we present a significant advancement that integrates entire Miller loop bit operations directly into the Schwartz-Zippel verification framework. While

the original approach focused on optimizing individual $\mathbb{F}_{q^{12}}$ multiplications through polynomial representation, our contribution expands this methodology to encompass complete Miller loop iterations.

Our key innovation lies in the development of three specialized verification patterns that handle different types of Miller loop operations: zero bits, non-zero positive/negative bits, and correction steps. Each pattern employs polynomial verification with carefully optimized commitment structures, reducing the constraint complexity while simultaneously decreasing auxiliary input requirements.

This integrated approach achieves substantial performance improvements in Miller loop computation while reducing calldata overhead — a critical consideration for blockchain deployment where calldata directly affects transaction costs. To validate these theoretical improvements, we systematically implemented and benchmarked each optimization technique.

1.5 Our Approach

We provide a comprehensive examination and systematic integration of these complementary state-of-the-art techniques, offering a consolidated analysis that quantifies their individual and cumulative performance benefits. Our methodology centers on implementing each optimization within a unified framework to measure the raw constraint savings and percentage improvements contributed by each technique when applied to a complete pairing verification pipeline.

These optimizations target different components of the pairing computation and can be implemented simultaneously: **El Housni’s** [Hou22] techniques optimize the overall constraint system structure and provide handcrafted, optimised formulae for operations throughout the pairing computation, **Feltroidprime’s** [Fel23] methods improve extension field arithmetic efficiency, and **Novakovic and Eagen’s** [NE24] approach eliminates final exponentiation overhead. Our systematic integration demonstrates how these techniques combine to achieve substantial overall performance gains, with each contributing measurable improvements to the total constraint count.

The original implementation of these techniques was developed in our Pairing library written in Cairo [Sri23]. The Garaga project [FE] on Starknet subsequently implemented these optimizations on top of Feltroidprime’s enhanced field multiplication techniques. Our contributions enabled Garaga to fit verifier’s auxiliary input within the transaction calldata constraints, while also making the Miller loop 50% more efficient. Through continued development by Feltroidprime, Rodrigo Ferreira, Liam Eagen, and other contributors, Garaga has evolved into one of the most performant solutions for zero-knowledge proof verification.

2 Background and Technical Roadmap

2.1 Pairing Computation

An elliptic curve pairing is a non-degenerate bilinear map $e : \mathbb{G}_1 \times \mathbb{G}_2 \rightarrow \mathbb{G}_T$, where $\mathbb{G}_1$ and $\mathbb{G}_2$ are groups of points on elliptic curves over finite fields, and $\mathbb{G}_T$ is a multiplicative group in an extension field. For practical cryptographic applications, we require that this map be efficiently computable.

The computation of an optimal ate pairing—the focus of our implementation—consists of two primary phases:

1. **Miller loop:** The core iterative computation that evaluates line functions at pairs of points from $\mathbb{G}_1$ and $\mathbb{G}_2$, accumulating intermediate results to produce a value in $\mathbb{F}_{q^{12}}$.
2. **Final Exponentiation:** This step raises the Miller loop output to the power $(q^{12} - 1)/r$, where r is the order of the groups, ensuring the result lies in the correct subgroup $\mathbb{G}_T$.

While both phases present optimization opportunities, the constraint-based verification of these computations introduces unique challenges and possibilities that differ significantly from native implementations.

2.2 Schwartz–Zippel lemma

The Schwartz–Zippel lemma (or more precisely, the Demillo–Lipton–Schwartz–Zippel lemma [DL78, Sch79, Zip79, Sch80]) provides a probabilistic method for testing polynomial identities. The lemma states that for a non-zero polynomial $P(x_1, x_2, \dots, x_n)$ of total degree d over a finite field $\mathbb{F}$, when evaluated at uniformly random field elements, the probability that P evaluates to zero is at most $d/|\mathbb{F}|$.

This probabilistic bound becomes negligible for cryptographically sized fields where $d \ll |\mathbb{F}|$, making it practical to verify polynomial identities by evaluation instead of coefficient-wise operations.

2.3 Roadmap

The following three sections systematically examine the optimization techniques that form the foundation of our work. Section 3 presents El Housni’s pioneering approach to pairings in rank-1 constraint systems [Hou22], establishing the theoretical framework for efficient pairing verification. Section 4 explores Feltroidprime’s enhanced extension field arithmetic [Fel23], which significantly improves field extension operation performance through polynomial commitment techniques. Section 5 details the residue witness approach of Novakovic and Eagen [NE24], which enables the elimination of final exponentiation overhead.

Each section provides theoretical background, implementation details, and quantitative performance analysis, building toward our novel contribution presented in Section 6. Our approach extends these foundations with optimized Miller loop operations that simultaneously reduce both computational complexity and verification calldata requirements.

3 Pairings in Rank-1 Constraint Systems

El Housni’s seminal work [Hou22] represents the first systematic exploration of pairing optimization within constraint-based computation models, providing techniques that extend beyond R1CS to benefit any arithmetized environment, including Starkware’s Cairo VM [Sta22].

3.1 Cost of inversions

Housni notes that the cost of inversions is much lower in constraint systems, almost the same as multiplication. Over $\mathbb{F}_p$, inversions cost 2 constraints: one for computing $x \cdot y$ where y is provided from the hint, and one equality check $x \cdot y \stackrel{?}{=} 1$. Division can be similarly reduced to just two constraints. For $c = a/b$ where c is the input from hint, we compute $b \cdot c$ and equality check $b \cdot c \stackrel{?}{=} a$. This is further generalised to field extensions.

For $x, y \in \mathbb{F}_{p^n}$
 $x/y = z$ (z as input from hint)
 Compute $t = y \cdot z$
 Equality check $t \stackrel{?}{=} x$
 Cost = 1 multiplication + n (n equality checks)

3.2 Miller loop

The paper’s section 6 is a comprehensive toolkit of Miller loop and final exponentiation optimizations, featuring handcrafted formulas and algorithmic insights specifically designed for constraint system efficiency. These optimizations form the foundation upon which we layer the subsequent advances discussed in this paper.

3.2.1 Affine coordinates

Since division cost about as much as multiplications, we can represent points in affine coordinates for lower operation counts. Point doubling and addition formulas are provided in section 6.1 of [Hou22]. Doubling costs 12C and addition costs 10C.

3.2.2 Avoid doubling

For double and add operations on non zero bits of the Miller loop, instead of computing $[2]S + Q$, computing $S + Q + S$ saves 2C per non zero bit. Taking this a step further, computation of y coordinate for the intermediate $S + Q$ point can be omitted and λ_1 can be directly used in computation of λ_2 . This **costs 17C**, that’s a **23% savings** over 22C for $[2]S + Q$.

3.2.3 Line evaluations

Following [ABLR14], the lines are sparse elements in $\mathbb{F}_{p^{12}}$ of the form $ay_P + bx_P \cdot w + c \cdot w^3$ with $a, b, c \in \mathbb{F}_{p^2}$. The multiplication of these sparse elements costs 39C. Housni proposes representing the lines in the form $1 + \frac{b'x_P}{y_P} \cdot w + \frac{c'}{y_P} \cdot w^3$ which makes it sparser (notice unit first coefficient resulting in free multiplication). This can be rewritten as,

$$1 + b' \cdot \frac{x_P}{y_P} \cdot w + c' \cdot \frac{1}{y_P} \cdot w^3$$

Terms $\frac{x_P}{y_P}$ and $\frac{1}{y_P}$ can be precomputed at the cost of 2C, and reused throughout the miller loop. Untwisting isomorphisms for this form are given on pages 15–16 in section 6.1 of [Hou22]. This representation only **costs 30C**, that’s a **30% savings** over 39C for the original representation.

3.3 Algebraic Torii

In section 5, El Housni explores Rubin and Silverberg’s torus-based arithmetic [RS03], which represents elements of $\mathbb{F}_{q^{2n}}$ using only n coordinates instead of $2n$, effectively compressing field extensions to half their size. While this compression necessitates field inversions for arithmetic operations—typically prohibitively expensive at $100\times$ the cost of multiplication in native implementations—the constraint system environment fundamentally alters this cost model. Inversions can be computed as hints outside the main constraint context and verified within the circuit using only a single multiplication constraint, making torus arithmetic highly attractive for reducing overall extension field operation complexity.

We will not go much deeper into this topic, but interested readers are encouraged to explore section 5 of El Housni’s paper [Hou22] for details.

3.4 Final exponentiation

In section 6.2 Final exponentiation of [Hou22], Housni presents a lot more tricks and optimizations for the final exponentiation step. Housni walks through easy part of the final exponentiation. And talks about when to apply which one of Granger-Scott algorithm [GS10] or SQR12345/SQR234 [Kar13] for squarings in the hard part. We will just walk through the costs of operations for comparison. Interested readers are encouraged to check out section 6.2 of El Housni’s paper [Hou22] for details on Final exponentiation.

3.5 Performance

Table 1 summarizes the arithmetic costs for $\mathbb{F}_{q^{12}}$ as presented in Housni’s paper.

Table 1: Arithmetic in Housni’s paper [Hou22]

	Mul	Square	Div	Sparse Mul
$\mathbb{F}_{q^{12}}$	54C	36C	66C	30C

4 Faster Extension Field Multiplications

Feltroidprime [Fel23] showcases a novel approach to extension field multiplications that exploits a counterintuitive property of circuit-based computation: operations traditionally viewed as expensive (such as polynomial divisions) can become highly efficient in constraint systems. Moving away from the conventional towered extension approach, he proposes returning to direct extensions of $\mathbb{F}_p$ to achieve significant performance improvements.

4.1 Isomorphisms among towered to direct extensions

Feltroidprime walks through the construction of mapping functions to go back and forth between direct and towered extensions.

He starts with the tower for BN254 (as used in gnark [Con]):

$$\mathbb{F}_{p^2}[u] = \mathbb{F}_p[u]/(u^2 + 1) \quad (1)$$

$$\mathbb{F}_{p^6}[v] = \mathbb{F}_{p^2}[v]/(v^3 - 9 - u) \quad (2)$$

$$\mathbb{F}_{p^{12}}[w] = \mathbb{F}_{p^6}[w]/(w^2 - v) \quad (3)$$

Obtains irreducible polynomials,

$$P_{12}(x) = x^{12} - 18x^6 + 82 \text{ and } P_6(x) = x^6 - 18v^3 + 82$$

And prepares isomorphisms between tower and direct representations [Fel] which we will not include in this paper for brevity.

4.2 Multiplication and Schwartz–Zippel lemma

Feltroidprime proposes using Euclidean division of polynomials to multiply $A, B \in \mathbb{F}_{p^{12}}$ or $A, B \in \mathbb{F}_{p^6}$. The computation happens as follows:

1. Treat A, B as polynomials of degree less than 11 or 5 respectively for $\mathbb{F}_{p^{12}}$ or $\mathbb{F}_{p^6}$ elements, with coefficients in $\mathbb{F}_p$.
2. Compute the product $C = A \cdot B$ as a polynomial of degree less than 22 for $\mathbb{F}_{p^{12}}$ or 10 for $\mathbb{F}_{p^6}$.
3. Perform Euclidean division of C by the irreducible polynomial $P_{12}(x)$ or $P_6(x)$ to obtain the quotient Q and remainder R :

$$\boxed{A(x) \cdot B(x) = Q(x) \cdot P_{12}(x) + R(x)} \text{ or } \boxed{A(x) \cdot B(x) = Q(x) \cdot P_6(x) + R(x)}$$

4. $R \in \mathbb{F}_{p^{12}}$ (or $R \in \mathbb{F}_{p^6}$) is used as the result of the multiplication between A and B .

Inside the circuit, this computation can be provided as a hint and verified with the Schwartz–Zippel lemma. The Fiat–Shamir heuristic can be used to generate the challenge with the transcript.

4.3 Deferred Evaluation and Batching

In Section 4, Feltroidprime discusses deferring the evaluation to the end of the algorithm and combining it with a polynomial commitment scheme to batch the verifications. This batching transforms hundreds of individual verifications into a single equation. So they can be checked at a single random point, amortizing the costs for Fiat-Shamir hashing.

The batching process is a random linear combination of all the equations to be verified. For n multiplications, we have:

$$\sum_{i=0}^{n-1} c_i * A_i(z) * B_i(z) = \sum_{i=0}^{n-1} c_i * R_i(x) + \sum_{i=0}^{n-1} c_i * Q_i(z) * P_{12}(z) \quad (4)$$

4.4 Performance

For performance analysis, we will use the following version of the batched verification outlined in 4.

$$\sum_{i=0}^{n-1} c_i * (A_i(z) * B_i(z) - R_i(x)) = \sum_{i=0}^{n-1} c_i * Q_i(z) * P_{12}(z)$$

The implementation receives the computed RHS polynomial aggregation $Q_{Agg} = \sum_{i=0}^{n-1} c_i * Q_i$ as hint. The LHS aggregation, $\sum_{i=0}^{n-1} c_i * A_i(z) * B_i(z) - R_i(x)$ is computed within the verifier circuit. While the amortized cost computing the $Q_{Agg}(z)$ and $P_{12}(z)$ once in the circuit approaches zero as the number of batched operations increases, we conservatively account for 1C per multiplication to provide fair cost comparison.

Multiplication costs 12C each for A , B and R evaluation, 1C for multiplying $A(z)$ and $B(z)$ evaluations, plus 1C for preparing the RLC c_i multiplication and conservative 1C for RHS aggregation.

For **squarings**, the evaluation of $A(z)$ is just used again, costing 12C less than multiplication.

Division costs multiplication plus additional 12C for assertion.

Sparse element multiplication costs Evaluations for A , B and R where B contains 5 non-zero coefficients, one out of which is unitary free multiplication. The costs comes at 31C which is 4C more than squaring which has free second polynomial because $A = B$ for squarings.

Table 2: Arithmetic in Feltroidprime’s paper [Fel23]

	Mul	Square	Div	Sparse Mul
$\mathbb{F}_{q^{12}}$	39C*	27C*	51C*	31C*
*Costs exclude Fiat-Shamir commitment overhead				

5 Eliminating the Final Exponentiation

The mathematical structure underlying pairing based cryptography involves equivalence classes. Tate pairings (as described in [NE24] and [Cos19]) result in the output in the quotient group $F_{q^k}/(F_{q^k})^r$. This requires the output of the Miller loop to be raised to a power of $(q^k - 1)/r$. This step, called **the final exponentiation**, normalizes the Miller loop output by clearing the r -th residue part. This normalization ensures equivalent pairing outputs produce identical results.

This operation is particularly expensive in extension fields like $\mathbb{F}_{q^{12}}$, because $(q^{12} - 1)/r$ becomes very large.

5.1 Equivalence class

Novakovic and Eagen [NE24] highlight that the final exponentiation can be eliminated when proving pairings. The expensive exponentiation can be replaced with an efficient residue check. The key insight is that instead of normalizing pairing outputs through

final exponentiation, we can directly verify equivalence using a residue-witness based approach.

Equivalence class representation: Two elements $A, B \in \mathbb{F}_{q^{12}}$ are equivalent in the pairing context if there exists some witness c such that $A \cdot c^r = B$.

$$A \equiv B \iff \exists c \in \mathbb{F}_{q^{12}} : A \cdot c^r = B$$

5.2 Equivalence via Residue Check

So we have,

$$A \cdot c^r \equiv A$$

After final exponentiation,

$$(A \cdot c^r)^{(q^{12}-1)/r} \equiv A^{(q^k-1)/r}$$

$$A^{(q^k-1)/r} \cdot c^{r \cdot (q^k-1)/r} \equiv A^{(q^k-1)/r}$$

$$A^{(q^k-1)/r} \cdot c^{q^k-1} \equiv A^{(q^k-1)/r}$$

With Fermat's Little Theorem,

$$A^{(q^k-1)/r} \cdot 1 \equiv A^{(q^k-1)/r}$$

If we expect the result of the pairing, $A \in \mathbb{F}_{q^k}$, to be equal to some desired output $B \in \mathbb{F}_{q^k}$, we can use the equivalence relation to establish there exists some $c \in \mathbb{F}_{q^k}$ such that $A \cdot c^r = B$.

This can be generalized to $A \cdot w^{rt} = B$ for $c = w^t$ in $A \cdot c^r = B$.

5.3 Finding optimal parameters

Section 4.2 and 4.3 of [NE24] dive into finding optimal parameters for the curve to establish:

$$A \cdot c^\lambda = B$$

5.3.1 Hasse bound and exponentiation by r

According to Hasse theorem [Sil09], $q + 1 - r \leq 2\sqrt{q}$ for an elliptic curve over finite field $\mathbb{F}_q$ and r the order of the curve.

The exponentiation by r can be written as exponentiation by $q+1-t$. Where $t \leq 2\sqrt{q}$ is much smaller exponentiation in worst case. And the computation of c^q is efficiently computed by exploiting the Frobenius endomorphisms.

5.3.2 Parameters for BN254 curve

Section 4.3 of [NE24] dives into finding optimal λ and computation of witness for Ethereum's BN254 of BN [BN05] curve family. For brevity, we will not dive into the details of the curve's construction.

$$t(x) = 6x^2 + 1$$

$$\lambda = 6x + 2 + q - q^2 + q^3$$

where x is the BN254 curve parameter

And also, $\lambda = 3rm'$, with r being the order of the pairing groups and m' the corresponding cofactor.

The exponentiation by λ decomposes as: $(6x + 2) + (q - q^2 + q^3)$

1. **Miller loop component** $(6x + 2)$ — Integrated into existing Miller loop squaring operations
2. **Frobenius component** $(q - q^2 + q^3)$ — Computed using efficient Frobenius mappings

This approach eliminates the computationally expensive final exponentiation entirely, replacing it with a few additional operations in the Miller loop and three efficient Frobenius mappings—a significant computational saving for constraint-based systems.

6 Miller loop Step as Polynomial

Having seen how final exponentiation can be eliminated through residue witnesses, we now turn to optimizing the Miller loop computation itself. Building upon Feltroidprime’s direct extension multiplication framework, we present a novel optimization that significantly improves Miller loop efficiency. The approach consolidates different operations in individual Miller loop steps into unified polynomial which is similarly committed and verified with Schwartz-Zippel, substantially reducing constraint count while maintaining cryptographic security.

6.1 Core Innovation

The Miller loop is comprised of steps for each individual bit, which in NAF (non-adjacent form) [HMOV04], may be +1, -1 or 0. Instead of performing a polynomial commitment for each individual extension field arithmetic operation separately, we propose the whole bit operation be viewed as one polynomial. And commitment be made to that polynomial.

To get a closer look into effectiveness of our scheme, we will make a real world application comparison for verification of three-pairing Miller loop which is common for verification of Groth16 [Gro16].

6.1.1 Comparison setup

1. Operate in a 3-pairing Miller loop setting.
2. Prepare the polynomial relationships to assert for a 0-bit operation.
3. We skip the evaluations and commitments of quotient polynomials ¹

¹The evaluation and commitment of quotient polynomials Q is skipped because they are aggregated into a single polynomial commitment at the end of the computation amortised across all operations.

Algorithm 1 Miller loop over Barreto–Naehrig curves [NE24]

Input: $P \in \mathbb{G}_1, Q \in \mathbb{G}_2$

Output: $f_{\lambda,Q}(P)$

```
1: Write  $s = 6t + 2 = \sum_{i=0}^L s_i 2^i$  where  $s_i \in \{-1, 0, 1\}$  and  $s_L = 1$ 
2:  $T \leftarrow Q, f \leftarrow 1$ 
3: for  $i = L - 2$  to  $0$  do
4:    $f \leftarrow f^2 \cdot l_{T,T}(P)$ 
5:    $T \leftarrow [2]T$ 
6:   if  $s_i = 1$  then
7:      $f \leftarrow f \cdot l_{T,Q}(P)$ 
8:      $T \leftarrow T + Q$ 
9:   end if
10:  if  $s_i = -1$  then
11:     $f \leftarrow f \cdot l_{T,-Q}(P)$ 
12:     $T \leftarrow T - Q$ 
13:  end if
14: end for
15:  $Q_1 \leftarrow \pi_p(Q), Q_2 \leftarrow \pi_p^2(Q)$ 
16:  $f \leftarrow f \cdot l_{T,Q_1}(P), T \leftarrow T + Q_1$ 
17:  $f \leftarrow f \cdot l_{T,-Q_2}(P), T \leftarrow T - Q_2$ 
18: return  $f$ 
```

4. Look at the Feltroidprime’s implementation [Fel23]

5. Compare it with our approach

6.1.2 Previous scheme

Adapting faster extension field multiplication [Fel23] idea to line 4 of Miller loop algorithm 1 for a 3–pair Miller loop, we have following polynomial relationships to assert.

$$\begin{aligned} F(x) \cdot F(x) &\stackrel{?}{=} Q_1(x) \cdot P_{12}(x) + T_1(x) \\ T_1(x) \cdot L_1(x) &\stackrel{?}{=} Q_2(x) \cdot P_{12}(x) + T_2(x) \\ T_2(x) \cdot L_2(x) &\stackrel{?}{=} Q_3(x) \cdot P_{12}(x) + T_3(x) \\ T_3(x) \cdot L_3(x) &\stackrel{?}{=} Q(x) \cdot P_{12}(x) + R(x) \end{aligned}$$

F, T_1, T_2, T_3, R cost 12C each for evaluation, 4C for $\mathbb{F}_q$ multiplication, and L_1, L_2, L_3 require 4C each for evaluation. All the remainder hints need to be provided for commitment, so this is 12 calldata commitment for R and intermediate remainder polynomials T_1, T_2, T_3 . **Costs sum up to 76C ($5 * 12C + 3 * 4C + 4C$) and 48 commitments** in $\mathbb{F}_q$ ($4 * 12$). Resulting remainder polynomial $R \in \mathbb{F}_{q^{12}}$ is the result of the Miller loop step.

6.1.3 Our scheme

In our approach, instead of working with individual arithmetic operations, we treat the entire bit operation as a single polynomial. In our scheme line 4 of 3–pair Miller loop

(Algorithm 1) looks like this.

$$F(x) \cdot F(x) \cdot L_1(x) \cdot L_2(x) \cdot L_3(x) \stackrel{?}{=} Q(x) \cdot P_{12}(x) + R(x)$$

Similar to the previous approach, F and R cost $12C$ each for evaluation, $4C$ for $\mathbb{F}_q$ multiplication, and L_1, L_2, L_3 require $4C$ each for evaluation, but we eliminate intermediate (T_* polynomials above) commitments and evaluations. Only the remainder hint needs commitment, so this is 12 calldata commitment for R . **Costs sum up to $40C$ ($2 * 12C + 3 * 4C + 4C$) and $12 \mathbb{F}_q$ commitments.**

This results in **47% fewer constraints** and a whopping **75% savings in the commitments**. When implemented on a blockchain for verification, calldata saving become far more critical.

6.2 Bit-Type Polynomials for 3-pairing

Our implementation uses different polynomials for different bits in the Miller loop. Let's walk through a 3-pairing setup which Inversions used widely for Groth16 verification. We will describe the polynomials for zero bits, non-zero bits and correction step.

6.2.1 Zero Bit Equation

For zero bit operations (line 4 with three $l_{T,T}(P)$ lines in Algorithm 1), the Miller loop just performs doubling. So for each pair we have Sparse_{01379} lines. The verification equation becomes:

$$F(x) \cdot F(x) \cdot L_1(x) \cdot L_2(x) \cdot L_3(x) = R(x) + Q(x) \cdot P_{12}(x)$$

Polynomial	Description	Degree	Coefficients
$F \in \mathbb{F}_{q^{12}}$	Miller loop accumulator	$11 \times 2 = 22$	12
$L_1, L_2, L_3 \in \text{Sparse}_{01379}$	The line functions	$9 \times 3 = 27$	$4 \times 3 = 12$
$R \in \mathbb{F}_{q^{12}}$	Remainder, The result	11	12
Total multiplications			36
Q Amortized	$\text{LHS}_{\text{Degree}} - R_{\text{Degree}}$	38	39

Table 3: Zero bit operation polynomial components

6.2.2 Non-Zero Bit Equation

For non-zero bit operations (line 4 and lines 6 to 12, with three pairs of lines in Algorithm 1), the Miller loop performs both doubling and addition steps. So for each pair we have two Sparse_{01379} lines. The verification equation becomes:

$$F(x) \cdot F(x) \cdot W(x) \cdot L_{d1}(x) \cdot L_{d2}(x) \cdot L_{d3}(x) \cdot L_{a1}(x) \cdot L_{a2}(x) \cdot L_{a3}(x) = R(x) + Q(x) \cdot P_{12}(x)$$

The residue witness W encapsulates the equivalence class relationship established by Novakovic and Eagen's final exponentiation elimination. The inverse of W is used for negative bits, this just changes the coefficients the equation remains the same.

Polynomial	Description	Degree	Coefficients
$F \in \mathbb{F}_{q^{12}}$	Miller loop accumulator	$11 \times 2 = 22$	12
$L_{d1}, L_{d2}, L_{d3} \in \text{Sparse}_{01379}$	Doubling line functions	$9 \times 3 = 27$	$4 \times 3 = 12$
$L_{a1}, L_{a2}, L_{a3} \in \text{Sparse}_{01379}$	Addition line functions	$9 \times 3 = 27$	$4 \times 3 = 12$
W or $W^{-1} \in \mathbb{F}_{q^{12}}$	Residue witness	11	12
$R \in \mathbb{F}_{q^{12}}$	Remainder, The result	11	12
Total multiplications			60
Q Amortized	$\text{LHS}_{\text{Degree}} - R_{\text{Degree}}$	76	77

Table 4: Non-zero bit operation polynomial components

6.2.3 Correction Step Equation

For the correction step operations, . The verification equation becomes:

$$F(x) \cdot F(x) \cdot L_1(x) \cdot L_2(x) \cdot L_3(x) = R(x) + Q(x) \cdot P_{12}(x)$$

Polynomial	Description	Degree	Coefficients
$F \in \mathbb{F}_{q^{12}}$	Miller loop accumulator	$11 \times 2 = 22$	12
$L_1, L_2, L_3 \in \text{Sparse}_{01379}$	The line functions	$9 \times 3 = 27$	$4 \times 3 = 12$
$R \in \mathbb{F}_{q^{12}}$	Remainder, The result	11	12
Total multiplications			36
Q Amortized	$\text{LHS}_{\text{Degree}} - R_{\text{Degree}}$	38	39

Table 5: Zero bit operation polynomial components

6.3 Polynomial Accumulation and implementation

The commitment can include all the required R s for each equation we want to prove. We arrive at the input elements through computation in the circuit, the commitments array is popped on reaching the point in the computation. We obtain an RLC (Random Linear Combination) of all LHS evaluations and precompute P_{12} . For example, for zero bits we go for,

$$F(x) \cdot F(x) \cdot L_1(x) \cdot L_2(x) \cdot L_3(x) - R(x) = Q(x) \cdot P_{12}(x)$$

or,

$$\text{LHS} = F(x) \cdot F(x) \cdot L_1(x) \cdot L_2(x) \cdot L_3(x) - R(x)$$

We then prepare an RLC of each equation to prove, for n equations we have,

$$\sum_{i=1}^n c_i \cdot \text{LHS}_i = \sum_{i=1}^n c_i \cdot Q_i(x) \cdot P_{12}(x)$$

Along with the commitments to each R we expect the prover to provide the correct RLC polynomial for $\sum_{i=1}^n c_i \cdot Q_i$. Here, x and all c_i are obtained from Fiat Shamir heuristics over all R_i commitments. This prevents the prover from guessing the randomness before commitment. But allows her to aggregate RLC for all quotients.

6.3.1 Prover

Prover performs the pairing verification herself, and computes and stores all R_i and Q_i . After the pairing computation, she uses same hashing method as the verifier to obtain all c_i . She then computes the RLC of all Q_i and sends it along with all R_i commitments.

6.3.2 Verifier

Verifier receives all the remainders, R_i , and RLC of quotients - $\sum_{i=1}^n c_i \cdot Q_i$. The verifier then prepares all c_i solely from the remainder coefficients. For the evaluation point x , she includes all coefficients from the quotient RLC in the Fiat-Shamir hash as well. This guarantees soundness in Fiat Shamir. Then she uses generated c_i for aggregating LHS _{i} evaluated at x . In the end, she subtracts the evaluation of quotients RLC at x times $P_{12}(x)$ from the aggregated LHS evaluation. If the result is zero, according to Schwartz–Zippel lemma, all the provided R_i are correct with a very high² probability.

6.4 Performance

Batching operations in Miller loop steps yielded substantial performance improvements. Processing all operations within a Miller loop step as a single polynomial and performing Schwartz-Zippel verification of polynomial multiplications significantly reduced the overall computational requirements of the pairing verification process. Combined with residue witness technique for final exponentiation elimination [NE24], these optimizations enabled processing of ZK proofs on Starknet even before the modulo builtin [Sta] became available, which later greatly improved finite field arithmetic operations. These techniques were subsequently integrated into Garaga, which remains the most efficient pairing implementation by operation count on Starknet to date.

References

- [ABLR14] Diego F. Aranha, Paulo S. L. M. Barreto, Patrick Longa, and Jefferson E. Ricardini. The realm of the pairings. In Tanja Lange, Kristin Lauter, and Petr Lisoněk, editors, *Selected Areas in Cryptography – SAC 2013*, pages 3–25, Berlin, Heidelberg, 2014. Springer Berlin Heidelberg.
- [BGW⁺22] Dan Boneh, Sergey Gorbunov, Riad S. Wahby, Hoeteck Wee, Christopher A. Wood, and Zhenfei Zhang. BLS Signatures. Internet-draft, Internet Engineering Task Force, June 2022. Work in Progress.
- [BN05] Paulo S. L. M. Barreto and Michael Naehrig. Pairing-friendly elliptic curves of prime order. Cryptology ePrint Archive, Paper 2005/133, 2005.
- [Con] Consensys. gnark: A fast zk-snark library. <https://github.com/Consensys/gnark>.
- [Cos19] Craig Costello. Pairings for beginners. Tutorial, 2019. Available at <https://www.craigcostello.com.au/tutorials>.

² $Pr[P(r_0, r_1, r_2, \dots, r_d) = 0] \leq \frac{d}{|\mathbb{F}|}$

- [DL78] Richard A. Demillo and Richard J. Lipton. A probabilistic remark on algebraic program testing. *Information Processing Letters*, 7(4):193–195, 1978.
- [FE] Feltroidprime and Keep Starknet Strange Exploration. Garaga: State-of-the-art elliptic curve tooling and snarks verification for cairo and starknet. <https://github.com/keep-starknet-strange/garaga>.
- [Fel] Feltroidprime. Bijection between fp12/fp6/fp2 tower representation and direct fp12/fp extension for bn254. <https://gist.github.com/feltroidprime/bd31ab8e0cbc0bf8cd952c8b8ed55bf5>.
- [Fel23] Feltroidprime. Faster extension field multiplications for emulated pairing circuits. <https://hackmd.io/@feltroidprime/B1eyHHXNT>, 2023. HackMD.
- [GGPR13] Rosario Gennaro, Craig Gentry, Bryan Parno, and Mariana Raykova. Quadratic span programs and succinct nizks without pcps. In *Advances in Cryptology - EUROCRYPT 2013*, pages 626–645. Springer, 2013. Foundational QAP framework for R1CS.
- [Gro16] Jens Groth. On the size of pairing-based non-interactive arguments. <https://eprint.iacr.org/2016/260>, 2016.
- [GS10] Robert Granger and Michael Scott. Faster squaring in the cyclotomic subgroup of sixth degree extensions. In Phong Q. Nguyen and David Pointcheval, editors, *Public Key Cryptography – PKC 2010*, pages 209–223, Berlin, Heidelberg, 2010. Springer Berlin Heidelberg.
- [GWC19] Ariel Gabizon, Zachary J. Williamson, and Oana Ciobotaru. Plonk: Permutations over lagrange-bases for oecumenical noninteractive arguments of knowledge. <https://eprint.iacr.org/2019/953>, 2019. Cryptology ePrint Archive, Paper 2019/953.
- [HMOV04] Darrel Hankerson, Alfred J. Menezes, and Scott Vanstone. *Guide to Elliptic Curve Cryptography*. Springer-Verlag, 2004. Non-adjacent form representation.
- [Hou22] Youssef El Housni. Pairings in rank-1 constraint systems. <https://eprint.iacr.org/2022/1162>, 2022.
- [Kar13] Koray Karabina. Squaring in cyclotomic subgroups. *Mathematics of Computation*, 82(281):555–579, 2013. Received October 22, 2010; revised August 22, 2011; published electronically June 27, 2012.
- [KZG10] Aniket Kate, Gregory M. Zaverucha, and Ian Goldberg. Constant-size commitments to polynomials and their applications. In Masayuki Abe, editor, *Advances in Cryptology - ASIACRYPT 2010*, pages 177–194, Berlin, Heidelberg, 2010. Springer Berlin Heidelberg.
- [NE24] Andrija Novakovic and Liam Eagen. On proving pairings. <https://eprint.iacr.org/2024/640>, 2024. Cryptology ePrint Archive, Paper 2024/640.

- [RS03] Karl Rubin and Alice Silverberg. Torus-based cryptography. In Dan Boneh, editor, *Advances in Cryptology - CRYPTO 2003*, pages 349–365, Berlin, Heidelberg, 2003. Springer Berlin Heidelberg.
- [Sch79] Jacob T. Schwartz. Probabilistic algorithms for verification of polynomial identities. In Edward W. Ng, editor, *Symbolic and Algebraic Computation*, pages 200–215, Berlin, Heidelberg, 1979. Springer Berlin Heidelberg.
- [Sch80] J. T. Schwartz. Fast probabilistic algorithms for verification of polynomial identities. *J. ACM*, 27(4):701–717, October 1980.
- [Sil09] Joseph H. Silverman. *The Arithmetic of Elliptic Curves*. Graduate Texts in Mathematics. Springer, New York, NY, 2 edition, 2009. 14 b/w illustrations.
- [Sri23] Shramee Srivastav. Cairo pairing. https://github.com/shramee/cairo_pairing, 2023.
- [Sta] Starknet. The mod builtin. <https://docs.starknet.io/build/corelib/core-circuit>. Cairo Book - Chapter 2.4.2.8.
- [Sta22] Starkware. Cairo: A turing-complete stark-friendly cpu architecture. <https://www.cairo-lang.org/>, 2022.
- [Zip79] Richard Zippel. Probabilistic algorithms for sparse polynomials. In Edward W. Ng, editor, *Symbolic and Algebraic Computation*, pages 216–226, Berlin, Heidelberg, 1979. Springer Berlin Heidelberg.