

Pairing Proofs and Polynomial Ring Toolkit

Anonymous Submission

Abstract. Elliptic curve pairings are fundamental to modern zero-knowledge proof systems, including Groth16, PLONK, and KZG polynomial commitments. In resource-constrained execution environments (Ethereum VM, BitVM, blockchain smart contracts) where computation is metered by gas fees, and in arithmetized computation systems (R1CS, Plonkish, Cairo AIR) where each computational step imposes significant prover overhead, the cost of pairing computation remains a practical bottleneck. We present a public-coin interactive proof system that consolidates the Miller loop iterations and final exponentiation into a single polynomial relation. The relation can then be verified as polynomial identities greatly reducing the pairing costs.

We present a full implementation in gnark [Cona]—the industry-standard framework for arithmetized circuit development—as a reusable polynomial ring toolkit in the **emulated** package. Benchmarked on a Groth16 verification simulation (2 fixed-point pairings + 1 full pairing + 1 $\mathbb{F}_{p,12}$ multiplication) on BN254, our gnark implementation achieves a **39% reduction in SCS constraints** ($1,890,771 \rightarrow 1,158,194$), a **60% reduction in commitments** ($739,198 \rightarrow 295,946$), and a **$\sim 42\%$ reduction** in compile-time memory and up to **52%** in solver memory. These improvements hold across both SCS and R1CS constraint systems and enable practical pairing verification in recursive proof systems and resource-constrained environments.

Keywords: Elliptic curve pairings · Zero-knowledge proofs · Polynomial identity testing · Blockchain verification · Recursive proofs.

1 Introduction

Blockchain technology continues to mature with improving scalability and infrastructure. However, privacy remains a significant barrier to entry for many users and applications. Zero-knowledge proofs (ZKPs) have emerged as a powerful solution to this challenge.

1.1 Motivation

Elliptic curve pairings represent fundamental building blocks for numerous cryptographic systems including Groth16[Gro16], PLONK[GWC19], BLS signature schemes[BGW⁺22], and KZG polynomial commitment schemes[KZG10].

However, **pairings are computationally intensive**. This creates critical bottlenecks in two important deployment contexts:

1. **Resource-constrained virtual machines** (Ethereum VM, BitVM, ZKVMs): Computation is metered — each operation incurs gas costs. A single pairing verification requiring thousands of field operations becomes economically prohibitive during peak network congestion.
2. **Arithmetized computation systems** (R1CS [GGPR13], Plonkish [GWC19], Cairo AIR [GPR21]): Each operation must be represented algebraically, adding significant prover overhead beyond native execution costs.

In such contexts, an efficient proof system for pairing correctness can replace expensive native execution inside the constrained environment.

1.2 Our Contribution

We present a public-coin interactive proof system for verifying pairing computations that consolidates verification into unified polynomial relationships under the Schwartz-Zippel lemma. By treating complete Miller loop iterations as single polynomial relationships rather than sequences of individual extension field operations [Fel23], our approach achieves:

- **39% reduction** in SCS constraints ($1,890,771 \rightarrow 1,158,194$ for BN254 Groth16 simulation)
- **60% reduction** in commitments ($739,198 \rightarrow 295,946$)
- **$\sim 42\%$ reduction** in compile-time memory and up to **52%** in solver memory

Our key insight is that instead of verifying individual field operations separately, we can verify entire Miller loop iterations as single polynomial identities. This consolidation dramatically reduces both the number of modular operations required and the communication overhead from intermediate commitments.

As a reusable artefact of independent interest, we also contribute a **polynomial ring toolkit** for gnark’s emulated arithmetic package. The toolkit packages the two-round interactive proof system into a clean API: a single `NewPolyRingCheck` call registers a modulus, `MulPolyRings` submits multiplications with deferred verification, and `performDeferredRingChecks` closes the circuit with two `api.Commit` challenges. Atop this, the `PolyRingAccumulator` provides a higher-level interface for building product chains: factors are enqueued via `Mul`, squarings via `Sqr`, and the accumulated product is evaluated lazily through `Eval`, with optional automatic degree-bounding to keep intermediate products manageable.

2 Background

2.1 Pairing Computation

An elliptic curve pairing is a non-degenerate bilinear map $e : \mathbb{G}_1 \times \mathbb{G}_2 \rightarrow \mathbb{G}_T$, where $\mathbb{G}_1$ and $\mathbb{G}_2$ are groups of points on elliptic curves over finite fields, and $\mathbb{G}_T$ is a multiplicative group in an extension field. For practical cryptographic applications, we require that this map be efficiently computable. We will focus on the optimal ate pairing [BGM⁺10] over Barreto-Naehrig (BN) curves, which is widely used due to its efficiency. We will specifically use the Ethereum’s BN254 curve, defined over a 254-bit prime field $\mathbb{F}_q$, with groups $\mathbb{G}_1, \mathbb{G}_2$ of order r and target group $\mathbb{G}_T$ contained in the 12th extension field $\mathbb{F}_{q^{12}}$.

The computation of an optimal ate pairing—the focus of our implementation—consists of two primary phases:

1. **Miller loop:** The core iterative computation that evaluates line functions at pairs of points from $\mathbb{G}_1$ and $\mathbb{G}_2$, accumulating intermediate results to produce a value in $\mathbb{F}_{q^{12}}$.
2. **Final Exponentiation:** This step raises the Miller loop output to the power $(q^{12} - 1)/r$, where r is the order of the groups, ensuring the result lies in the correct subgroup $\mathbb{G}_T$.

While both phases present optimization opportunities, the constraint-based verification of these computations introduces unique challenges and possibilities that differ significantly from native implementations.

Algorithm 1 Optimal Ate pairing over Barreto-Naehrig curves [BGM⁺10]**Input:** $P \in \mathbb{G}_1, Q \in \mathbb{G}_2$ **Output:** $a_{opt}(Q, P)$

```

1: Write  $s = 6t + 2 = \sum_{i=0}^{L-1} s_i 2^i$  where  $s_i \in \{-1, 0, 1\}$  and  $s_L = 1$ 
2:  $T \leftarrow Q, f \leftarrow 1$ 
3: for  $i = L - 2$  to  $0$  do
4:    $f \leftarrow f^2 \cdot l_{T,T}(P)$ 
5:    $T \leftarrow [2]T$ 
6:   if  $s_i = 1$  then
7:      $f \leftarrow f \cdot l_{T,Q}(P)$ 
8:      $T \leftarrow T + Q$ 
9:   end if
10:  if  $s_i = -1$  then
11:     $f \leftarrow f \cdot l_{T,-Q}(P)$ 
12:     $T \leftarrow T - Q$ 
13:  end if
14: end for
15:  $Q_1 \leftarrow \pi_p(Q), Q_2 \leftarrow \pi_p^2(Q)$ 
16:  $f \leftarrow f \cdot l_{T,Q_1}(P), T \leftarrow T + Q_1$ 
17:  $f \leftarrow f \cdot l_{T,-Q_2}(P), T \leftarrow T - Q_2$ 
18:  $f \leftarrow f^{(p^{12}-1)/r}$ 
19: return  $f$ 

```

84 **2.2 Schwartz-Zippel lemma**

85 The Schwartz-Zippel lemma (or more precisely, the Demillo-Lipton-Schwartz-Zippel
 86 lemma [DL78, Sch79, Zip79, Sch80]) provides a probabilistic method for testing poly-
 87 nomial identities. The lemma states that for a non-zero polynomial $P(x_1, x_2, \dots, x_n)$ of
 88 total degree d over a finite field $\mathbb{F}$, when evaluated at uniformly random field elements, the
 89 probability that P evaluates to zero is at most $d/|\mathbb{F}|$.

90 This probabilistic bound becomes negligible for cryptographically sized fields where
 91 $d \ll |\mathbb{F}|$, making it practical to verify polynomial identities by evaluation instead of
 92 coefficient-wise operations.

93 **2.3 Computation Models**

94 We implement and benchmark our toolkit in gnark [Cona], the industry-standard Go library
 95 for arithmetized circuit development, targeting both SCS (Sparse Constraint System, the
 96 native Plonkish backend) and R1CS constraint systems. We map each field multiplication
 97 to a constraint (denoted by **C**) to maintain compatibility with the R1CS model used in
 98 established literature. In R1CS [Gro16], prover complexity is dominated by multiplication
 99 gates while additions are free. To synchronize our theoretical analysis with the benchmarks
 100 provided by El Housni [Hou23], we define our cost metric based on field multiplications.

101 **2.3.1 R1CS Cost Model**

102 We measure computational cost in terms of *constraints* (denoted by **C**), representing the
 103 fundamental unit of work in Rank-1 constraint systems. One constraint corresponds to
 104 one multiplication gate.

105 **Basic Field Operations in $\mathbb{F}_q$:**

- 106 • **Addition/Subtraction:** Free

- **Multiplication:** One constraint (1C)
- **Division/Inversion:** One multiplication + one equality check

Commitment Overhead: Beyond constraint costs, we account for *calldata* or *commitment overhead*—the number of field elements the prover must commit to (via Fiat-Shamir hashing in non-interactive proofs). For blockchain deployments, this overhead directly impacts transaction costs and is often the dominant factor.

2.3.2 Cost Model Comparison

Operations that are traditionally expensive (inversions, divisions) become cheap in constraint systems when computed via witnesses. This inverts traditional optimization priorities: we can afford inversions to save multiplications, enabling affine coordinate systems and other techniques impractical in native implementations.

Recognizing and exploiting these unique computational properties is essential for unlocking hidden performance potential. Youssef El Housni’s pioneering work **Pairings in Rank-1 Constraint Systems** [Hou23] represents, to our knowledge, the first systematic exploration of these optimization opportunities, providing a breakthrough in the practicality of recursive proving by thoughtfully leveraging the distinctive cost model inherent to constraint-based systems.

3 Related Work

The landscape of efficient pairing verification in constraint systems has evolved through three key contributions that our work directly builds upon.

1. El Housni’s **Pairings in Rank-1 Constraint Systems** [Hou23] established the theoretical framework for optimizing pairings in arithmetized environments, introducing efficient Miller loop formulas in affine coordinates and sparse line function representations. This work inverts the usual cost model: inversions become cheap (2C) while multiplications remain the dominant cost.
2. Feltroidprime [Fel23] introduced polynomial identity testing for $\mathbb{F}_{q^{12}}$ arithmetic via the Schwartz-Zippel lemma. Each multiplication $A \cdot B$ is expressed as $A(x) \cdot B(x) = Q(x) \cdot P_{12}(x) + R(x)$ and verified at a single random point, with the quotients Q batched via random linear combination. Our work applies the same polynomial framework but at a coarser granularity: entire Miller loop steps rather than individual multiplications, eliminating intermediate remainder commitments entirely.
3. Novakovic and Eagen’s **On Proving Pairings** [NE24] showed that final exponentiation can be replaced by a residue witness check, reducing final exponentiation cost by over 85% for BN254. We integrate their residue witness c directly into our polynomial step equations (Section 4.2.2), unifying the Miller loop and final exponentiation into the same polynomial verification framework.

4 Pairing Operations as Polynomials

Algorithm 2 gives the pairing computation combining El Housni’s R1CS optimizations [Hou23] with Novakovic-Eagen’s residue witness technique [NE24]. The 27th root of unity W scales the Miller loop output f to be a cubic residue (see [NE24] §4.3); if f is not already a cubic residue, multiplying by W^w , $w \in \{1, 2\}$ corrects this.

Our contribution is to express each step of this algorithm as a single multi-operand polynomial product verified via the Schwartz-Zippel lemma, rather than as a sequence of individual binary multiplications as in [Fel23].

Algorithm 2 Optimal Ate pairing with Residue witness check

Input: $P \in \mathbb{G}_1, Q \in \mathbb{G}_2, c \in \mathbb{F}_{q^{12}}, w \in \{0, 1, 2\}$
 Internal $W, W^2 \in \mathbb{F}_{q^{12}}$ ▷ W is 27th root of unity

Output: $a_{opt}(Q, P)$

- 1: Write $s = 6t + 2 = \sum_{i=0}^{L-1} s_i 2^i$ where $s_i \in \{-1, 0, 1\}$ and $s_L = 1$
- 2: $c^{-1} \leftarrow 1/c$ ▷ Inverse Residue witness
- 3: $f \leftarrow c^{-1}$ ▷ Initialize with inverse residue witness
- 4: $T \leftarrow Q$ ▷ T_i for each pair for multi-pairing
- Every operation involving T, P or Q is repeated for T_i, P_i and Q_i for multi-pairing.

- 5: **for** $i = L - 2$ to 0 **do**
- 6: $f \leftarrow f^2$
- 7: $f \leftarrow f \cdot l_{T,T}(P)$ ▷ Repeat for T_i, P_i
- 8: $T \leftarrow [2]T$ ▷ Repeat for T_i
- 9: **if** $s_i = 1$ **then**
- 10: $f \leftarrow f \cdot c^{-1}$ ▷ Residue witness
- 11: $f \leftarrow f \cdot l_{T,Q}(P)$ ▷ Repeat for T_i, Q_i, P_i
- 12: $T \leftarrow T + Q$ ▷ Repeat for T_i, Q_i
- 13: **else if** $s_i = -1$ **then**
- 14: $f \leftarrow f \cdot c$ ▷ Residue witness
- 15: $f \leftarrow f \cdot l_{T,-Q}(P)$ ▷ Repeat for $T_i, -Q_i, P_i$
- 16: $T \leftarrow T - Q$ ▷ Repeat for $T_i, -Q_i$
- 17: **end if**
- 18: **end for**

- 19: $Q_1 \leftarrow \pi_p(Q), Q_2 \leftarrow \pi_p^2(Q)$ ▷ Repeat for Q_{1i}, T_i
- 20: $f \leftarrow f \cdot l_{T,Q_1}(P)$ ▷ Repeat for Q_{1i}, T_i, P_i
- 21: $T \leftarrow T + Q_1$ ▷ Repeat for Q_{1i}, T_i
- 22: $f \leftarrow f \cdot l_{T,-Q_2(Q)}(P)$ ▷ Repeat for Q_i, T_i, P_i
- 23: **if** $w = 1$ **then**
- 24: $f \leftarrow f \cdot W$
- 25: **else if** $w = 2$ **then**
- 26: $f \leftarrow f \cdot W^2$
- 27: **end if**
- 28: $f \leftarrow f \cdot c^{-q} \cdot c^{q^2} \cdot c^{-q^3}$ ▷ Uses Frobenius maps, Negative exponents use c^{-1}

- 29: **return** f

4.1 Comparative Analysis

We compare the two schemes for a 3-pair Miller loop zero-bit step (the most common case in BN254’s NAF[DHM04] representation of $s = 6t + 2$). Here $F \in \mathbb{F}_{q^{12}}$ is the accumulator, $L_i \in \text{Sparse}_{01379}$ are the three doubling line functions, $R \in \mathbb{F}_{q^{12}}$ is the updated accumulator, and Q is the quotient polynomial. Sparse_{01379} denotes the subspace of $\mathbb{F}_{q^{12}}$ elements whose only nonzero coefficients are at degrees 0, 1, 3, 7, 9—the form that BN254 line functions naturally take [Hou23], reducing their evaluation cost from 12C to 4C.

4.1.1 Previous scheme

Adapting [Fel23] to lines 6–7 of Algorithm 2 for a 3-pair Miller loop gives:

$$F(x) \cdot F(x) \stackrel{?}{=} Q_1(x) \cdot P_{12}(x) + T_1(x)$$

$$T_1(x) \cdot L_1(x) \stackrel{?}{=} Q_2(x) \cdot P_{12}(x) + T_2(x)$$

$$T_2(x) \cdot L_2(x) \stackrel{?}{=} Q_3(x) \cdot P_{12}(x) + T_3(x)$$

$$T_3(x) \cdot L_3(x) \stackrel{?}{=} Q(x) \cdot P_{12}(x) + R(x)$$

Evaluating F , T_1 – T_3 , R costs 12C each and L_1 – L_3 cost 4C each; all four remainder polynomials require commitment. **Total: 76C and 48 $\mathbb{F}_q$ commitments** (4×12).

4.1.2 Our scheme

Our scheme consolidates lines 6–7 into a single polynomial identity:

$$F(x) \cdot F(x) \cdot L_1(x) \cdot L_2(x) \cdot L_3(x) \stackrel{?}{=} Q(x) \cdot P_{12}(x) + R(x)$$

By eliminating intermediate T_* polynomials, only R requires commitment. **Total: 41C and 12 $\mathbb{F}_q$ commitments** ($2 \times 12C + 3 \times 4C + 5C$).

This results in **46% fewer constraints** and a **75% reduction in commitments**. When verifying ZK proofs on-chain, commitment savings are often more critical than constraint savings, due to the calldata costs associated with transmitting polynomial hints.

4.1.3 Performance Comparison

Table 1 compares the computational costs and commitment overhead between the schemes for a single 0-bit Miller loop operation with 3 pairs.

Table 1: 3-Pair Miller loop 0-bit constraints comparison

Scheme	Constraints	$\mathbb{F}_q$ Commitments	Improvement
[Hou23]	132C	—	Baseline
[Fel23]	76C*	48	42%
Our scheme	41C*	12	69%
[Hou23]: 1 square (36C) + 3 sparse mul ($3 \times 30C$) = 132C [Fel23]: 5 evals ($5 \times 12C$) + 3 line evals ($3 \times 4C$) + 4 muls ($4 \times 1C$) = 76C Our scheme: 2 evals ($2 \times 12C$) + 3 line evals ($3 \times 4C$) + 5C (muls + RLC) = 41C *Costs exclude Fiat-Shamir commitment overhead			

4.2 Operation Polynomials

Our implementation uses different polynomials for different bits in the Miller loop. Let us walk through a 3-pairing setup which Inversions used widely for Groth16 verification. We will describe the polynomials for zero bits, non-zero bits and correction step.

4.2.1 Miller loop: Zero Bit Equation

As seen in 4.1.2 for zero bit operations (line 6 and 7 with three $l_{T,T}(P)$ lines in Algorithm 2), the verification equation is:

$$F^2(x) \cdot L_1(x) \cdot L_2(x) \cdot L_3(x) = R(x) + Q(x) \cdot P_{12}(x) \quad (1)$$

Table 2: Zero bit operation polynomial components

Polynomial	Description	Degree	Coefficients
$F \in \mathbb{F}_{q^{12}}$	Miller loop accumulator	$11 \times 2 = 22$	12
$L_1, L_2, L_3 \in \text{Sparse}_{01379}$	The line functions	$9 \times 3 = 27$	$4 \times 3 = 12$
$R \in \mathbb{F}_{q^{12}}$	Remainder, The result	11	12
Total constraints	36C Evals, 4C Muls and 1C RLC		41
Q (<i>Amortized See 5.1</i>)	$\deg(\text{LHS}) - \deg(P_{12}) + 1$	38	39

4.2.2 Miller loop: Non-Zero Bit Equation

For non-zero bit operations (lines 10, 11, 14, 15 with three pairs of lines in Algorithm 2), the Miller loop performs both doubling and addition steps. So for each pair we have two Sparse_{01379} lines. The verification equation becomes:

$$F^2(x) \cdot W(x) \cdot L_{d1}(x) \cdot L_{d2}(x) \cdot L_{d3}(x) \cdot L_{a1}(x) \cdot L_{a2}(x) \cdot L_{a3}(x) = R(x) + Q(x) \cdot P_{12}(x) \quad (2)$$

Table 3: Non-zero bit operation polynomial components

Polynomial	Description	Degree	Coefficients
$F \in \mathbb{F}_{q^{12}}$	Miller loop accumulator	$11 \times 2 = 22$	12
$L_{d1}, L_{d2}, L_{d3} \in \text{Sparse}_{01379}$	Doubling line functions	$9 \times 3 = 27$	$4 \times 3 = 12$
$L_{a1}, L_{a2}, L_{a3} \in \text{Sparse}_{01379}$	Addition line functions	$9 \times 3 = 27$	$4 \times 3 = 12$
W or $W^{-1} \in \mathbb{F}_{q^{12}}$	Residue witness	11	12
$R \in \mathbb{F}_{q^{12}}$	Remainder, The result	11	12
Total constraints	60C Evals, 8C Muls and 1C RLC		69
Q (<i>Amortized See 5.1</i>)	$\deg(\text{LHS}) - \deg(P_{12}) + 1$	76	77

The residue witness W encapsulates the equivalence class relationship established by Novakovic and Eagen’s final exponentiation elimination. The inverse of W is used for negative bits, this just changes the coefficients the equation remains the same.

4.2.3 Miller loop: Frobenius mapping Equation

The correction step finalizes the Miller loop by incorporating the Frobenius endomorphism corrections (lines 20 and 22 in Algorithm 2). Again, we have two Sparse_{01379} lines for each pair, for π_p and $-\pi_p^2$ mappings, without any squaring of the accumulator. The verification equation becomes:

$$F(x) \cdot L_{1\pi}(x) \cdot L_{2\pi}(x) \cdot L_{3\pi}(x) \cdot L_{1\pi^2}(x) \cdot L_{2\pi^2}(x) \cdot L_{3\pi^2}(x) = R(x) + Q(x) \cdot P_{12}(x) \quad (3)$$

Table 4: Correction step polynomial components

Polynomial	Description	Degree	Coefficients
$F \in \mathbb{F}_{q^{12}}$	Miller loop accumulator	11	12
$L_{1\pi}, L_{2\pi}, L_{3\pi} \in \text{Sparse}_{01379}$	π_p correction lines	$9 \times 3 = 27$	$4 \times 3 = 12$
$L_{1\pi^2}, L_{2\pi^2}, L_{3\pi^2} \in \text{Sparse}_{01379}$	$-\pi_p^2$ correction lines	$9 \times 3 = 27$	$4 \times 3 = 12$
$R \in \mathbb{F}_{q^{12}}$	Remainder, The result	11	12
Total constraints	48C Evals, 6C Muls and 1C RLC		55
Q (<i>Amortized See 5.1</i>)	$\deg(\text{LHS}) - \deg(P_{12}) + 1$	54	55

4.2.4 Final Exponentiation Equation

The final exponentiation check is handled as described in [NE24]. The equation handles the cubic scale factor multiplication, and **Frobenius component** ($q - q^2 + q^3$) of the final exponentiation (lines 24, 26 and 28 in Algorithm 2). The verification equation looks like:

$$F(x) \cdot c^{-q}(x) \cdot c^{q^2}(x) \cdot c^{-q^3}(x) \cdot W^w(x) = 1 + Q(x) \cdot P_{12}(x) \quad (4)$$

where $w \in \{0, 1, 2\}$ determines the cubic root of unity scaling factor, and the Frobenius mappings are computed using the residue witness c and its inverse and provided as input. Cubic root is sparse element with just 2 coefficients the degree ranging from 8 to 10 because of positioning of the coefficients. Our R is now 1 for a successful proof verification.

Table 5: Final exponentiation polynomial components

Polynomial	Description	Degree	Coefficients
$F \in \mathbb{F}_{q^{12}}$	Final Miller loop result	11	12
$c^{-q} \in \mathbb{F}_{q^{12}}$	Residue witness Frobenius $-q$	11	12
$c^{q^2} \in \mathbb{F}_{q^{12}}$	Residue witness Frobenius q^2	11	12
$c^{-q^3} \in \mathbb{F}_{q^{12}}$	Residue witness Frobenius $-q^3$	11	12
$W^w \in \mathbb{F}_{q^{12}}$	Cubic scale factor	10	2
$R = 1$	Unity (constant)	0	1
Total constraints	51C Evals, 4C Muls and 1C RLC		56
Q (<i>Amortized See 5.1</i>)	$\deg(\text{LHS}) - \deg(P_{12}) + 1$	43	44

With this we have all the polynomials we need for verifying entire pairing operation. Next we will look into high level overview of pairing verification via an interactive oracle proof.

4.3 Security

Each verification equation takes the form $P = \prod_i A_i(x) - Q(x) \cdot P_{12}(x) - R(x)$ where $\deg(P) < 2^8$.¹ By the Schwartz-Zippel lemma (Section 2.2):

Soundness: $\delta_s \leq d/|\mathbb{F}_q| < 2^8/2^{254} = 2^{-246}$ for BN254.

Completeness: An honest prover always produces $P(x) = 0$ identically, so the verifier always accepts. Euclidean division guarantees unique Q, R with $\deg(R) < \deg(P_{12}) = 12$.

¹For k -pair pairings, $\deg(P) \approx 88 + 18(k - 3)$, well below 2^8 for all practical k . This bound can be extended to 2^{10} with negligible soundness impact.

5 Full Pairing Proof with Quotient Batching

We now combine the polynomial equations from Section 4.2 into a complete two-round public-coin interactive proof, further reducing verification cost by batching all quotients into a single accumulated polynomial.

5.1 Batching Polynomial Identity Testing

All n polynomial equations from Section 4.2 are batched into a single verification. The prover commits to remainders $\{R_i\}_{i=0}^{n-1}$ and receives challenge z . Using z , the prover forms the accumulated quotient $Q_{\text{acc}} = \sum_{i=0}^{n-1} z^i \cdot Q_i$. The verifier then samples a second challenge point c and checks:

$$\sum_{i=0}^{n-1} z^i \cdot [A_i(c) \cdot B_i(c) \cdots X_i(c) - R_i(c)] = Q_{\text{acc}}(c) \cdot P_{12}(c) \quad (5)$$

This reduces n polynomial verifications (degrees 38–76, see Tables 2–5) to a single evaluation of a degree-76 polynomial.

5.2 Prover's Algorithms

The prover computes quotients Q_i and remainders R_i for each equation in Section 4.2. For each step, it calls POLYMULTIDIV (Algorithm 3), which multiplies the input polynomials and performs standard Euclidean division by the irreducible P_{12} . After receiving challenge z , the prover forms $Q_{\text{acc}} = \sum_i z^i Q_i$ and sends $\mathcal{R} = \{R_i\}$ and Q_{acc} to the verifier. The full prover procedure is given in Algorithm 4.

Algorithm 3 PolyMulDiv: Multiply polynomials and divide by P_{12}

Input: $P_1, \dots, P_n \in \mathbb{F}_q[x]$; Divisor $P_{12} \in \mathbb{F}_q[x]$

Output: Q, R s.t. $\prod P_i = Q \cdot P_{12} + R$, $\deg(R) < \deg(P_{12})$

- 1: $R \leftarrow \prod_{i=1}^n P_i$ $\triangleright$ Polynomial product
 - 2: Perform Euclidean division of R by P_{12} to get (Q, R)
 - 3: **return** (Q, R)
-

5.3 Verifier's Algorithms

The verifier receives $P \in \mathbb{G}_1$, $Q \in \mathbb{G}_2$, residue witness c [NE24], the W (27th root of unity), the remainder array $\mathcal{R}$, and Q_{acc} from the prover. Using challenge z from the first round and a freshly sampled x , it mirrors the prover's pairing loop while evaluating each polynomial at x via Algorithm 5, accumulating the RLC e_{acc} . The final check is Equation 5. The full verifier is given in Algorithm 6.

Algorithm 4 Pairing Prover: Preparing polynomials

Input: $P \in \mathbb{G}_1, Q \in \mathbb{G}_2, c \in \mathbb{F}_{q^{12}}, w \in \{0, 1, 2\}$
 Internal: $W, W^2 \in \mathbb{F}_{q^{12}}$ $\triangleright W$ is 27th root of unity

Output: Arrays $\mathcal{Q}, \mathcal{R}$ of quotients and remainders
 Write $s = 6t + 2 = \sum_{i=0}^{L-1} s_i 2^i$ where $s_i \in \{-1, 0, 1\}$ and $s_L = 1$

- 1: $c^{-1} \leftarrow 1/c$ $\triangleright$ Inverse Residue witness
- 2: $f \leftarrow c^{-1}$ $\triangleright$ Initialize with inverse residue witness
- 3: $T \leftarrow Q$ $\triangleright T_i$ for each pair for multi-pairing

Every operation involving T, P or Q is repeated for T_i, P_i and Q_i for multi-pairing.

Miller loop bits

- 4: **for** $i = L - 2$ to 0 **do**
- 5: $L_{\text{dbl}} \leftarrow l_{T,T}(P), T \leftarrow [2]T$
- 6: **if** $s_i = 0$ **then**
- 7: $Q, R \leftarrow \text{PolyMulDiv}(f, f, L_{\text{dbl}})$ $\triangleright$ Algorithm 3, add L_{*i} for each P_i, Q_i
- 8: **else if** $s_i = 1$ **then**
- 9: $L_{\text{add}} \leftarrow l_{T,Q}(P), T \leftarrow T + Q$
- 10: $Q, R \leftarrow \text{PolyMulDiv}(f, f, c^{-1}, L_{\text{dbl}}, L_{\text{add}})$ $\triangleright$ Algorithm 3, add L_{*i} for each P_i, Q_i
- 11: **else if** $s_i = -1$ **then**
- 12: $L_{\text{add}} \leftarrow l_{T,-Q}(P), T \leftarrow T - Q$
- 13: $Q, R \leftarrow \text{PolyMulDiv}(f, f, c, L_{\text{dbl}}, L_{\text{add}})$ $\triangleright$ Algorithm 3, add L_{*i} for each P_i, Q_i
- 14: **end if**
- 15: $f \leftarrow R, \mathcal{Q}.\text{append}(Q), \mathcal{R}.\text{append}(R)$
- 16: **end for**

Miller loop Frobenius correction

- 17: $Q_1 \leftarrow \pi_p(Q), Q_2 \leftarrow \pi_p^2(Q)$
- 18: $L_\pi \leftarrow l_{T,Q_1}(P)$
- 19: $T \leftarrow T + Q_1$
- 20: $L_{-\pi^2} \leftarrow l_{T,-Q_2}(P)$
- 21: $Q, R \leftarrow \text{PolyMulDiv}(f, L_\pi, L_{-\pi^2})$ $\triangleright$ Algorithm 3, add L_{*i} for each P_i, Q_i
- 22: $f \leftarrow R, \mathcal{Q}.\text{append}(Q), \mathcal{R}.\text{append}(R)$

Final exponentiation: Cubic scaling + Frobenius mappings

- 23: **if** $w = 0$ **then**
- 24: $Q, R \leftarrow \text{PolyMulDiv}(f, c^{-q}, c^{q^2}, c^{-q^3})$ $\triangleright$ Algorithm 3
- 25: **else if** $w = 1$ **then**
- 26: $Q, R \leftarrow \text{PolyMulDiv}(f, W, c^{-q}, c^{q^2}, c^{-q^3})$ $\triangleright$ Algorithm 3
- 27: **else if** $w = 2$ **then**
- 28: $Q, R \leftarrow \text{PolyMulDiv}(f, W^2, c^{-q}, c^{q^2}, c^{-q^3})$ $\triangleright$ Algorithm 3
- 29: **end if**
- 30: $\mathcal{Q}.\text{append}(Q), \mathcal{R}.\text{append}(R)$ $\triangleright$ Final $\mathcal{Q}$ and $\mathcal{R}$ accumulation, skip f
- 31: **return** $\mathcal{Q}, \mathcal{R}$

Algorithm 5 EvalPoly: Evaluate polynomial product minus remainder via Horner’s method

Input: Polynomials $P_1, \dots, P_n, R \in \mathbb{F}_q[x]$ of degree $\leq d$, evaluation point $z \in \mathbb{F}_q$


Output: $e \in \mathbb{F}_q$ where $e = \prod_{i=1}^n P_i(z) - R(z)$

```

1: function Eval ( $P = (p_0, p_1, \dots, p_d), z$ ): return (evaluation  $\in \mathbb{F}_q$ )
2:   Uses Horner’s method, body not included for brevity
3:  $e \leftarrow \text{EVAL}(P_1, z)$  ▷ Evaluate first polynomial
4: for  $i = 2$  to  $n$  do
5:    $e \leftarrow e \cdot \text{EVAL}(P_i, z)$  ▷ Multiply in each subsequent evaluation
6: end for
7:  $e \leftarrow e - \text{EVAL}(R, z)$  ▷ Subtract remainder evaluation
8: return  $e$ 

```

5.4 Interactive Proof



5.5 Fiat-Shamir Transformation

Applying the Fiat-Shamir transformation [FS87] with random oracle H gives a non-interactive proof with knowledge soundness $\delta_s + 3(m^2 + 1)2^{-\lambda}$ [BSCS16], where $m = 2$ (hash queries) and $\lambda = 254$ (BN254 hash output length).

Gnark’s multi-commitment mechanism. In gnark, each `api.Commit` call registers committed variables in the constraint system; the two sequential calls in our protocol produce the two challenges z and x via gnark’s multi-round commitment scheme [BSB, BSB23].

Algorithm 6 Pairing Verifier: Evaluating remainders**Input:** $P \in \mathbb{G}_1, Q \in \mathbb{G}_2, c \in \mathbb{F}_{q^{12}}, w \in \{0, 1, 2\}, \mathcal{R}$ array of remainders, $Q_{\text{acc}} \in \mathbb{F}_q[x]$ Internal: 27th root of unity, $W, W^2 \in \mathbb{F}_{q^{12}}$, first round challenge, $z \xleftarrow{\$} \mathbb{F}_q$ **Output:** $\text{pairing}(P, Q) \stackrel{?}{=} 1$ Write $s = 6t + 2 = \sum_{i=0}^{L-1} s_i 2^i$ where $s_i \in \{-1, 0, 1\}$ and $s_L = 1$

1: $c^{-1} \leftarrow 1/c$ ▷ Inverse Residue witness
2: $f \leftarrow c^{-1}$ ▷ Initialize with inverse residue witness
3: $T \leftarrow Q$ ▷ T_i for each pair for multi-pairing
4: $x \xleftarrow{\$} \mathbb{F}_q$ ▷ Sample evaluation point from field
5: $e_{\text{acc}} \leftarrow 0, a_{\text{rlc}} \leftarrow 1$ ▷ Evaluation accumulator and random linear combination multiplier
Every operation involving T, P or Q is repeated for T_i, P_i and Q_i for multi-pairing.

Miller loop bits

6: **for** $i = L - 2$ to 0 **do**
7: $L_{\text{dbl}} \leftarrow l_{T,T}(P), T \leftarrow [2]T$
8: $R \leftarrow \mathcal{R}.\text{pop_front}()$ ▷ Fetch R from prover
9: **if** $s_i = 0$ **then**
10: $e_{\text{acc}} \leftarrow e_{\text{acc}} + a_{\text{rlc}} * \text{EvalPoly}(R, x, f, f, L_{\text{dbl}})$ ▷ Algo 5, add L_{*i} for each P_i, Q_i
11: **else if** $s_i = 1$ **then**
12: $L_{\text{add}} \leftarrow l_{T,Q}(P), T \leftarrow T + Q$
13: $e_{\text{acc}} \leftarrow e_{\text{acc}} + a_{\text{rlc}} * \text{EvalPoly}(R, x, f, f, c^{-1}, L_{\text{dbl}}, L_{\text{add}})$ ▷ Algo 5, add L_{*i} for each P_i, Q_i
14: **else if** $s_i = -1$ **then**
15: $L_{\text{add}} \leftarrow l_{T,-Q}(P), T \leftarrow T - Q$
16: $e_{\text{acc}} \leftarrow e_{\text{acc}} + a_{\text{rlc}} * \text{EvalPoly}(R, x, f, f, c, L_{\text{dbl}}, L_{\text{add}})$ ▷ Algo 5, add L_{*i} for each P_i, Q_i
17: **end if**
18: $f \leftarrow R, a_{\text{rlc}} \leftarrow a_{\text{rlc}} \cdot z$ ▷ Update random linear combination factor
19: **end for**

Miller loop Frobenius correction


20: $R \leftarrow \mathcal{R}.\text{pop_front}()$ ▷ Fetch R from prover
21: $Q_1 \leftarrow \pi_p(Q), Q_2 \leftarrow \pi_p^2(Q)$
22: $L_\pi \leftarrow l_{T,Q_1}(P)$
23: $T \leftarrow T + Q_1$
24: $L_{-\pi^2} \leftarrow l_{T,-Q_2}(P)$
25: $e_{\text{acc}} \leftarrow e_{\text{acc}} + a_{\text{rlc}} * \text{EvalPoly}(R, x, f, L_\pi, L_{-\pi^2})$ ▷ Algo 5, add L_{*i} for each P_i, Q_i
26: $f \leftarrow R, a_{\text{rlc}} \leftarrow a_{\text{rlc}} \cdot z$ ▷ Update random linear combination factor

Final exponentiation: Cubic scaling + Frobenius mappings

27: $R \leftarrow \mathcal{R}.\text{pop_front}()$ ▷ Fetch R from prover
28: **if** $w = 0$ **then**
29: $e_{\text{acc}} \leftarrow e_{\text{acc}} + a_{\text{rlc}} * \text{EvalPoly}(R, x, f, c^{-q}, c^{q^2}, c^{-q^3})$ ▷ Algo 5
30: **else if** $w = 1$ **then**
31: $e_{\text{acc}} \leftarrow e_{\text{acc}} + a_{\text{rlc}} * \text{EvalPoly}(R, x, f, W, c^{-q}, c^{q^2}, c^{-q^3})$ ▷ Algo 5
32: **else if** $w = 2$ **then**
33: $e_{\text{acc}} \leftarrow e_{\text{acc}} + a_{\text{rlc}} * \text{EvalPoly}(R, x, f, W^2, c^{-q}, c^{q^2}, c^{-q^3})$ ▷ Algo 5
34: **end if**

35: **return** $Q_{\text{acc}}(x) \stackrel{?}{=} e_{\text{acc}} \wedge R \stackrel{?}{=} 1$ ▷ Polynomial identity and expected pairing output check

5.6 Non-Interactive Proof



Concrete commitment costs: In gnark, each call to `api.Commit` registers committed variables in the constraint system; the precise prover overhead depends on the backend (Groth16 or Plonk) [BSB, BSB23].

5.7 Security Analysis

Throughout this section, $z, x \in \mathbb{F}_q$ denote field element challenges (scalars), while polynomials are elements of $\mathbb{F}_q[x]$ with x as the formal indeterminate. In particular, z^j is a scalar coefficient, not a polynomial term.

5.7.1 Algebraic Intuition

Before the formal soundness proof it is instructive to consider why an incorrect remainder is difficult to hide. Suppose a malicious prover corrupts exactly one remainder: $R'_j = R_j + \Delta_j$ with $\Delta_j \in \mathbb{F}_q[x]$, $\Delta_j \neq 0$, $\deg(\Delta_j) \leq 11$. Since z is obtained *after* the commitment to $\mathcal{R}$, the prover cannot adjust any other R_k to compensate for Δ_j . To successfully deceive the verifier, the prover must find a valid forged quotient $Q'_{\text{acc}} \in \mathbb{F}_q[x]$ satisfying

$$\sum_{i=0}^{n-1} z^i (A_i - R'_i) = Q'_{\text{acc}} \cdot P_{12}$$

All $R'_i = R_i$ except for $R'_j = R_j + \Delta_j$. Substituting $R'_j = R_j + \Delta_j$ gives

$$Q'_{\text{acc}} \cdot P_{12} = Q_{\text{acc}} \cdot P_{12} - z^j \Delta_j(x)$$

$$Q'_{\text{acc}} = Q_{\text{acc}} - \frac{z^j \Delta_j(x)}{P_{12}(x)}.$$

Here $z^j \in \mathbb{F}_q$ is a nonzero scalar (with overwhelming probability over the choice of z) and $\Delta_j(x) \in \mathbb{F}_q[x]$ is a nonzero polynomial of degree at most 11. Since P_{12} is irreducible of degree 12 and $\deg(z^j \Delta_j) \leq 11 < 12 = \deg(P_{12})$, the polynomial P_{12} does not divide $z^j \Delta_j$ in $\mathbb{F}_q[x]$; hence $z^j \Delta_j / P_{12}$ is not a polynomial. Consequently no valid $Q'_{\text{acc}} \in \mathbb{F}_q[x]$ can satisfy the equation for this single-corruption case.

The same degree argument extends immediately to an adversary who corrupts *multiple* remainders: for any scalar z , the weighted sum $\sum_i z^i \Delta_i(x)$ is still a polynomial in x of degree at most 11 (since each z^i is a scalar and each $\deg(\Delta_i) \leq 11$), so it cannot be divisible by P_{12} of degree 12 unless it is the zero polynomial, which requires every $\Delta_i = 0$. The formal soundness proof below re-derives this uniformly via Schwartz-Zippel.

5.7.2 Soundness

Theorem 1 (Soundness). *Let n be the number of polynomial equations batched via a random linear combination into the bivariate polynomial identity:*

$$P_{\text{final}}(z, x) = \sum_{i=0}^{n-1} z^i \cdot \left[\prod_j P_{i,j}(x) - R_i(x) \right] - Q_{\text{acc}}(x) \cdot P_{12}(x) = 0$$

Let the maximum degree in x of the polynomial products be $d_x < 2^8$. For a prover committing to incorrect remainders, the probability of successful forgery δ_s over randomly chosen $(z, x) \in \mathbb{F}_q^2$ is bounded by:

$$\delta_s \leq \frac{n - 1 + d_x}{|\mathbb{F}_q|}$$

Proof. If the prover commits to an incorrect remainder $R_i(x)$, the evaluation yields a non-zero error polynomial $P_{\text{err}}(z, x)$. The structure of $P_{\text{final}}(z, x)$ establishes the following degree bounds for $P_{\text{err}}(z, x)$:

- **Degree in z :** at most $n - 1$ (due to the z^i weighting).
- **Degree in x :** at most $d_x < 2^8$ (due to the polynomial products).

By the Schwartz-Zippel lemma, the probability that the non-zero polynomial P_{err} evaluates to zero is bounded by the total degree divided by the field size:

$$\Pr[P_{\text{err}}(z, x) = 0] \leq \frac{\deg_z(P_{\text{err}}) + \deg_x(P_{\text{err}})}{|\mathbb{F}_q|} \leq \frac{n - 1 + d_x}{|\mathbb{F}_q|}$$

286

□

Concrete parameters for BN254: With $n = 67$ equations and $d_x \leq 256$:

$$\delta_s \leq \frac{66 + 256}{2^{254}} < \frac{2^9}{2^{254}} = 2^{-245}.$$

Non-interactive proof via Fiat-Shamir : From Section 5.5, under the random oracle model the knowledge soundness error becomes [BSCS16]:

$$\delta'_s = \delta_s + 3(m^2 + 1) \cdot 2^{-\lambda}$$

where $m = 2$ (hash queries) and $\lambda = 254$ (hash output length), giving

$$\delta'_s \leq 2^{-245} + 15 \cdot 2^{-254} < 2^{-244}.$$

287 *Remark 1* (Correctness of polynomial evaluation). Theorem 1 assumes that the verifier
 288 evaluates all polynomials at the challenge point x without arithmetic error. The circuit
 289 implementation uses the deferred-reduction inner-product optimisation of Section 6.5
 290 (Section 6.5.1), which accumulates limb-wise products over $\mathbb{Z}$ before a single final reduction.
 291 Correctness of this optimisation requires that no intermediate limb value reaches the
 292 native field characteristic p . We verify this bound concretely for BN254 in Lemma 1
 293 (Section 6.5.1).

294 5.7.3 Completeness

295 For an honest prover executing Algorithm 4 with valid inputs:

1. **Polynomial construction:** At each operation, the prover computes following polynomials:

$$\prod_j P_{i,j} = Q_i \cdot P_{12} + R_i$$

2. **RLC formation:** Given challenge z , the prover forms:

$$Q_{\text{acc}} = \sum_{i=0}^{n-1} z^i \cdot Q_i$$

- 296 3. **Verifier's check:** The verifier evaluates at point x :

$$\begin{aligned} \text{LHS} &= \sum_{i=0}^{n-1} z^i \cdot \left[\prod_j P_{i,j}(x) - R_i(x) \right] \\ &= \sum_{i=0}^{n-1} z^i \cdot Q_i(x) \cdot P_{12}(x) \quad (\text{by construction}) \\ &= P_{12}(x) \cdot \sum_{i=0}^{n-1} z^i \cdot Q_i(x) \\ &= P_{12}(x) \cdot Q_{\text{acc}}(x) = \text{RHS} \end{aligned}$$

301 Since the equality holds for any evaluation point x , the protocol satisfies perfect
 302 completeness.

303 6 Polynomial Ring Toolkit

304 The protocol of Section 5 is packaged as a reusable toolkit in the `emulated` package of
 305 gnark [Conb], parameterised by any emulated field $\mathbb{F}_q$ (via `FieldParams` type T) and any
 306 modulus polynomial $P \in \mathbb{F}_q[x]$. The toolkit operates in a recursive proof setting where
 307 the prover's computations are performed outside the circuit via `hints` (Section 6.2). And
 308 verifier's checks are implemented as constraints inside the circuit.

309 Refer to our fork at <https://anonymous.4open.science/r/tiny-gnark-DDCD> for
 310 code.

311 6.1 Ring Representation

312 A ring element is essentially a coefficient slice `[]*Element[T]` in ascending degree order. A
 313 `nil` entry is treated as zero by `InnerProductNoReduce`—the term is skipped entirely—so
 314 sparse polynomials carry no wasted constraint variables.

315 The optional `modEvalFn` of type `EvalFnType[T] = ([*Element[T]]) → *Element[T]`
 316 allows handling small-integer coefficients efficiently by multiplying by the small positive
 317 integer first and negating afterwards, rather than multiplying by the full-width field
 318 representative $q - c$.

319 **Example: BN254 $\mathbb{F}_{p^{12}}$**

320 $P_{12}(x) = x^{12} - 18x^6 + 82$ has nine zero coefficients (stored as `nil`) and two small non-zero
 321 ones. The `modEvalFn` computes $P_{12}(\alpha) = \text{MulConst}(\alpha^0, 82) + \text{Neg}(\text{MulConst}(\alpha^6, 18)) + \alpha^{12}$,
 322 keeping intermediate values small instead of multiplying by $q - 18$.

323 6.2 Hint Functions

324 Hints allow defining computations outside of a circuit. A hint defines an annotated
 325 function in the circuit at compile time. Inputs and outputs are injected at the solving
 326 time. We define these two hints to implement the prover's steps in the two-round protocol
 327 of Section 5.4.

328 **polyRingMulHint** implements `POLYMULDIV` (Algorithm 3): it multiplies the input
 329 polynomials, divides by P to obtain (Q, R) , and serialises the result as `nbQTerms` | Q limbs |
 330 `nbRemTerms` | R limbs. Inputs are packed as `nbBits` | `nbLimbs` | `nbPoly` | q | inputs | P ,
 331 where each polynomial is prefixed by its term count.

332 **quotientsRLCHint** computes the accumulated quotient $Q_{\text{acc}}[j] = \sum_i z^i \cdot Q_i[j] \bmod q$
 333 coefficient-wise, corresponding to the prover's step in Algorithm 4.

334 6.3 API

335 The toolkit exposes three entry points. `NewPolyRingCheck` registers a ring group (mod-
 336 ulus P plus an optional small-coefficient evaluator) and returns a handle shared by every
 337 multiplication under that modulus; multiple groups with different moduli may coexist.
 338 `MulPolyRings` computes $R = (\prod_j A_j) \bmod P$ outside the circuit via `polyRingMulHint`,
 339 range-checks R 's limbs inside, and enqueues the check $(A_j; R; Q)$ in the group; for longer
 340 product chains a `PolyRingAccumulator` wrapper automatically triggers intermediate re-
 341 ductions when the accumulated degree exceeds a configurable bound. At circuit finalisation,
 342 `performDeferredRingChecks` implements Equation (5) via two sequential `api.Commit`
 343 calls [BSB] (producing challenges z and x) and one `AssertIsEqual` per group, so no
 344 in-circuit work beyond remainder range-checks is performed until this point. Full API
 345 signatures and a step-by-step code walkthrough are available in the repository.²

347 6.4 Usage

348 Core API

349 Register the ring group once at initialisation:

```
350 fp, _ := emulated.NewField[emulated.BN254Fp](api)
351 modPoly := fp.MakePoly([]any{82, 0, 0, 0, 0, 0, -18, 0, 0, 0, 0, 0, 1})
352 ring := fp.NewPolyRingCheck(modPoly, nil)
353 // Each Miller-loop step then calls:
354 rem, _ = f.MulPolyRings(ring, f, f, L1, L2, L3)
```

²<https://anonymous.4open.science/r/tiny-gnark-DDCD>

355 and `performDeferredRingChecks` is invoked once by the gnark backend at finalisation,
 356 running the full two-round protocol of Section 5.

357 **Accumulate a product chain**

```
358 acc := f.NewPolyRingAccumulator(group, targetDeg)
359 acc.Mul(A)
360 acc.Mul(B)
361 acc.Sqr()      // enqueues (A·B)2 without multiplying yet
362 rem := acc.Eval()
```

363 **6.5 Implementation Details and Performance**

364 **6.5.1 Schoolbook vs Horner’s Method for Polynomial Evaluation**

Horner’s method for evaluations across emulated elements causes a linear increase in the limbs count per coefficient. This necessitates reductions after each multiplication. Instead, we cache powers of the evaluation point x in $\mathcal{X} = (x^0, x^1, \dots, x^d)$ and evaluate each polynomial as an inner product of $\mathcal{X}$ and the coefficient array $\mathbf{F}$:

$$\mathcal{X} \cdot \mathbf{F} = \mathcal{X}_0 \cdot F_0 + \mathcal{X}_1 \cdot F_1 + \dots + \mathcal{X}_d \cdot F_d$$

365 In the emulated field context, EVAL in Algorithm 5 is replaced by this inner product
 366 with $\mathcal{X}$. This allows us to use the inner product optimization described below.

367 **6.5.2 Native field operations for inner products**

368 Modular reductions can be deferred or entirely elided for operations expressible as inner
 369 products, such as polynomial evaluations and random linear combinations. This optimiza-
 370 tion leverages the capacity headroom provided by representing small emulated limbs (e.g.,
 371 64-bit limbs) within a significantly larger native field element.

372 **Deferred Reduction:** Intermediate scalars are permitted to grow over the integers $\mathbb{Z}$
 373 during limb-wise multiplications and linear accumulations.

374 **Preservation of Identity:** Algebraic identities are strictly preserved provided the
 375 maximum accumulated limb magnitude remains bounded below the native characteristic
 376 p .

377 **Handling Asymmetry:** In circuits with asymmetric multiplicative depths, operands
 378 may accumulate as disparate polynomial multiples of the emulated characteristic p_{emul} .
 379 Consequently, a terminal single-element reduction suffices for equivalence testing, bypassing
 380 the constraint overhead of intermediate normalization.

381 **Lemma 1** (Overflow safety for deferred reduction). *Let $\mathbb{F}_q$ be emulated with limbs of width*
 382 *b bits inside a native field $\mathbb{F}_p$. Consider evaluating a degree- d polynomial via an inner*
 383 *product over the cached power vector $\mathcal{X}$. Since limbs are accumulated independently, each*
 384 *result-limb is a sum of $d + 1$ products of two b -bit values, so its magnitude is bounded by*
 385 *$\beta = 2^{2b + \lceil \log_2(d+1) \rceil}$.*

386 *The deferred single-element reduction is safe whenever $\beta < \log_2 p$.*

387 **Corollary 1** (BN254 pairing instantiation). *For BN254 pairing verification: $b = 64$,*
 388 *$d \leq 256$, and $p \approx 2^{254}$. Then*

$$389 \quad \beta = 2 \cdot 64 + \lceil \log_2 257 \rceil = 128 + 9 = 137 < 254,$$

390 *so the deferred reduction is safe.*

Table 6: Performance improvements in Garaga: Cairo implementation on Starknet

MULMOD		ADDMOD		COMMITTS		VM STEPS	
Before	After	Before	After	Before	After	Before	After
BN254: 3-Pairing with Miller loop and final exponentiation							
19,142	12,656	21,686	15,142	5,689	2,807	338,678	212,902
33.9%		30.2%		50.7%		37.1%	
BLS12-381: 3-Pairing with Miller loop and final exponentiation							
16,484	11,287	20,669	15,420	5,421	3,167	325,757	219,155
31.5%		25.4%		41.6%		32.7%	
MULMOD/ADDMOD: Field operations in $\mathbb{F}_p$ via built-in function calls							
COMMITTS: Use transaction calldata and Poseidon hash operations for Fiat-Shamir							
VM STEPS: CairoVM Steps [Sta22] (CPU Cycles)							

7 Conclusion

We have shown that choosing the right granularity for polynomial identity testing—entire Miller loop steps rather than individual field multiplications—yields substantial and measurable improvements in arithmetized pairing verification. The key technique, collapsing a composite multi-operand product into a single Euclidean division check with batched quotient accumulation, is simple to implement and broadly applicable to other multi-step cryptographic computations in arithmetized environments.

7.1 Applications and Impact

Our approach delivers maximum impact in resource-constrained execution environments where custom auxiliary data can be provided for cheaper verification. **Recursive proof systems** benefit the most greatly reducing outer circuit complexity from fewer commitments and operations. **ZKVMs** (RISC Zero, SP1, Jolt) benefit naturally as they express computation as constraint systems compatible with our polynomial framework. **Calldata-constrained blockchains** (Ethereum L2s, BitVM) also benefit, as fewer commitments means smaller proof transcripts and associated calldata costs.

This is already implemented in-production in Garaga [FE] for CairoVM and is currently the standard for all pairing based ZK proofs on Starknet. The polynomial ring toolkit 6 in gnark, submitted as a part of this work, will directly impact `evmprecompiles/ecpair` precompile via `AssertMillerLoopAndFinalExpIsOne`.

7.2 Implementation and Performance

The Garaga library [FE] provided an earlier proof of concept for this technique in Cairo/Starknet (2024). The resulting reductions in commitment overhead and modulo operations (Table 6) were sufficient to fit Groth16 proof verification within Starknet’s transaction limits for the first time. However, Cairo [Sta22] is a specialized environment with its own cost metrics, making direct comparison with mainstream ZK frameworks difficult. This work situates the same technique within gnark—the standard framework for field emulation in arithmetized circuits—providing a systematic benchmark against an industry baseline.

Table 7 presents benchmark results from our gnark implementation of the polynomial ring toolkit, measured on a Groth16 verification simulation on BN254. The end-to-end results confirm both directions of saving: **39% SCS constraint reduction** (1,890,771 $\rightarrow$ 1,158,194) and **60% commitment reduction** (739,198 $\rightarrow$ 295,946), with strong secondary improvements across compile-time memory ($\sim 42\text{--}43\%$) and solver memory ($\sim 26\%$ SCS, $\sim 52\%$ R1CS).

Table 7: Performance comparison: gnark implementation, Groth16 simulation on BN254 (2 fixed-point pairings + 1 full pairing + 1 $\mathbb{F}_{p^{12}}$ mul)

Metric	gnark Baseline	This Work	Reduction
SCS (Plonkish) Constraint System			
Nb Constraints	1,890,771	1,158,194	38.7%
Nb Commitments	739,198	295,946	60.0%
Compile Time	775.64 ms	406.71 ms	47.6%
Compile Memory	1.77 GB	1.03 GB	41.8%
Solve Time	479.61 ms	356.42 ms	25.7%
Solve Memory	510.36 MB	376.57 MB	26.2%
R1CS Constraint System			
Nb Constraints	588,013	353,881	39.8%
Nb Commitments	739,198	295,946	60.0%
Compile Time	780.40 ms	423.09 ms	45.8%
Compile Memory	2.09 GB	1.19 GB	43.1%
Solve Time	320.84 ms	214.73 ms	33.1%
Solve Memory	330.54 MB	158.51 MB	52.0%

Commitments: Fiat-Shamir `api.Commit` calls. Hardware: Apple M3 Pro (`goarch:arm64`).

Circuit benchmarked: https://archive.org/download/ppprt_groth16_simulation_test

Our implementation: <https://anonymous.4open.science/r/tiny-gnark-DDCD/>

Benchmark circuit. Both implementations are evaluated against the same gnark circuit³,

Step 1: A 2-pair fixed- $\mathbb{G}_2$ pairing check $e(P_1, Q_1) \cdot e(P_2, Q_2) = 1$, where both $\mathbb{G}_2$ points use precomputed line evaluations (`PairingCheck` with fixed Q).

Step 2: A single in-circuit Miller loop $ML(P_3, Q_3)$ combined with a Miller loop result f computed *outside* the circuit, followed by a final-exponentiation-equals-one assertion (`AssertMillerLoopAndFinalExpIsOne`).

This structure simulates constraints for a real Groth16 verification: the two fixed- $\mathbb{G}_2$ pairings, one full pairing, and one $\mathbb{F}_{p^{12}}$ multiplication. These recreate the four pairing operations with fixed trusted setup points.

The gnark baseline numbers were obtained by running the benchmark function against the upstream gnark `emulated` package (github.com/consensys/gnark). The **This Work** numbers were obtained from the same benchmark file compiled against our fork, which replaces the per-multiplication deferred checks in the `emulated` package with the polynomial ring toolkit described in this paper. No changes to the circuit definition or test harness were required; only the underlying `emulated` arithmetic backend differs between the two runs.

References

- [BGM⁺10] Jean-Luc Beuchat, Jorge Enrique González-Díaz, Shigeo Mitsunari, Eiji Okamoto, Francisco Rodríguez-Henríquez, and Tadanori Teruya. High-speed software implementation of the optimal ate pairing over barreto-naehrig curves. In *Pairing-Based Cryptography – Pairing 2010*, volume 6487 of *Lecture Notes in Computer Science*, pages 21–39. Springer, 2010.
- [BGW⁺22] Dan Boneh, Sergey Gorbunov, Riad S. Wahby, Hoeteck Wee, Christopher A. Wood, and Zhenfei Zhang. BLS Signatures. Internet-draft, Internet Engineering Task Force, June 2022. Work in Progress.

³Circuit definition: https://archive.org/download/ppprt_groth16_simulation_test

- [BSB] Alexandre Belling, Azam Soleimanian, and Olivier Bégassat. Multi-round, multi-commitment: A proposal to extend gnark’s `api.Commit` to support multi-round protocols. <https://hackmd.io/x8KsadW3RRyX7YTCFJikHg>.
- [BSB23] Alexandre Belling, Azam Soleimanian, and Olivier Bégassat. Recursion over public-coin interactive proof systems; faster hash verification. In *Proceedings of the 2023 ACM SIGSAC Conference on Computer and Communications Security, CCS ’23*, pages 1422–1436, New York, NY, USA, 2023. Association for Computing Machinery.
- [BSCS16] Eli Ben-Sasson, Alessandro Chiesa, and Nicholas Spooner. Interactive oracle proofs. In *Theory of Cryptography – TCC 2016-B*, volume 9986 of *Lecture Notes in Computer Science*, pages 31–60. Springer, 2016.
- [Cona] Consensys. gnark: A fast zk-snark library. <https://github.com/Consensys/gnark>.
- [Conb] Consensys. gnark emulated fields. <https://github.com/Consensys/gnark/tree/master/std/algebra/emulated>.
- [DHM04] Scott Vanstone Darrel Hankerson and Alfred Menezes. *Guide to Elliptic Curve Cryptography*. Springer-Verlag, 2004. Non-adjacent form representation.
- [DL78] Richard A. Demillo and Richard J. Lipton. A probabilistic remark on algebraic program testing. *Information Processing Letters*, 7(4):193–195, 1978.
- [FE] Feltroidprime and Keep Starknet Strange Exploration. Garaga: State-of-the-art elliptic curve tooling and snarks verification for cairo and starknet. <https://github.com/keep-starknet-strange/garaga>.
- [Fel23] Feltroidprime. Faster extension field multiplications for emulated pairing circuits. <https://hackmd.io/@feltroidprime/BleyHHXNT>, 2023. HackMD.
- [FS87] Amos Fiat and Adi Shamir. How to prove yourself: Practical solutions to identification and signature problems. In Andrew M. Odlyzko, editor, *Advances in Cryptology – CRYPTO’ 86*, pages 186–194, Berlin, Heidelberg, 1987. Springer Berlin Heidelberg.
- [GGPR13] Rosario Gennaro, Craig Gentry, Bryan Parno, and Mariana Raykova. Quadratic span programs and succinct nizks without pcps. In *Advances in Cryptology – EUROCRYPT 2013*, pages 626–645. Springer, 2013. Foundational QAP framework for R1CS.
- [GPR21] Lior Goldberg, Shahar Papini, and Michael Riabzev. Cairo – a turing-complete STARK-friendly CPU architecture. Cryptology ePrint Archive, Paper 2021/1063, 2021.
- [Gro16] Jens Groth. On the size of pairing-based non-interactive arguments. In *Advances in Cryptology – EUROCRYPT 2016*, volume 9666 of *Lecture Notes in Computer Science*, pages 305–326. Springer, 2016.
- [GWC19] Ariel Gabizon, Zachary J. Williamson, and Oana Ciobotaru. Plonk: Permutations over lagrange-bases for oecumenical noninteractive arguments of knowledge. <https://eprint.iacr.org/2019/953>, 2019. Cryptology ePrint Archive, Paper 2019/953.

- 493 [Hou23] Youssef El Housni. Pairings in rank-1 constraint systems. In *Topics in*
494 *Cryptology – CT-RSA 2023*, volume 13871 of *Lecture Notes in Computer*
495 *Science*, pages 321–345. Springer, 2023.
- 496 [KZG10] Aniket Kate, Gregory M. Zaverucha, and Ian Goldberg. Constant-size com-
497 mitments to polynomials and their applications. In Masayuki Abe, editor,
498 *Advances in Cryptology - ASIACRYPT 2010*, pages 177–194, Berlin, Heidelberg,
499 2010. Springer Berlin Heidelberg.
- 500 [NE24] Andrija Novakovic and Liam Eagen. On proving pairings. [https://eprint.](https://eprint.iacr.org/2024/640)
501 [iacr.org/2024/640](https://eprint.iacr.org/2024/640), 2024. Cryptology ePrint Archive, Paper 2024/640.
- 502 [Sch79] Jacob T. Schwartz. Probabilistic algorithms for verification of polynomial
503 identities. In Edward W. Ng, editor, *Symbolic and Algebraic Computation*,
504 pages 200–215, Berlin, Heidelberg, 1979. Springer Berlin Heidelberg.
- 505 [Sch80] J. T. Schwartz. Fast probabilistic algorithms for verification of polynomial
506 identities. *J. ACM*, 27(4):701–717, October 1980.
- 507 [Sta22] Starkware. Cairo: A turing-complete stark-friendly cpu architecture. [https:](https://www.cairo-lang.org/)
508 [//www.cairo-lang.org/](https://www.cairo-lang.org/), 2022.
- 509 [Zip79] Richard Zippel. Probabilistic algorithms for sparse polynomials. In Edward W.
510 Ng, editor, *Symbolic and Algebraic Computation*, pages 216–226, Berlin, Hei-
511 delberg, 1979. Springer Berlin Heidelberg.